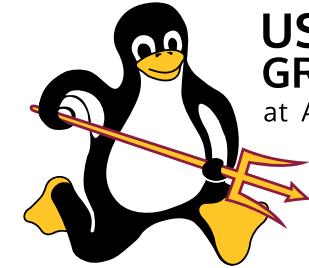


Qubes OS

A Technical Introduction

Daylam Tayari



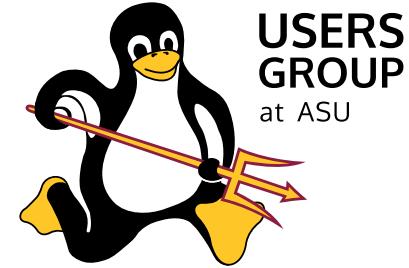
Technical?

Don't worry – This won't be exclusively technical

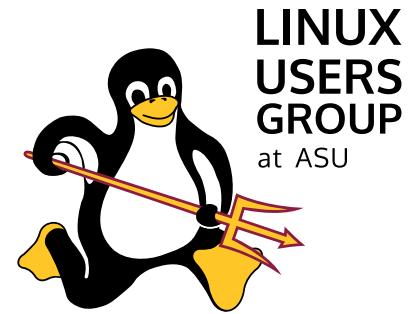
Overview → Specificities → Technical Details

Questions at any point?
Ask!

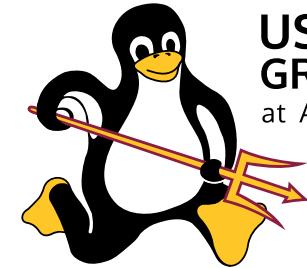
Talk Layout



1. Overview
2. Usage and Security Implementations
3. Drivers, Interfaces and Policies
4. Technical Analysis of the Architecture
5. Security and Attack Surfaces



Section 1 – Overview



What is Qubes?

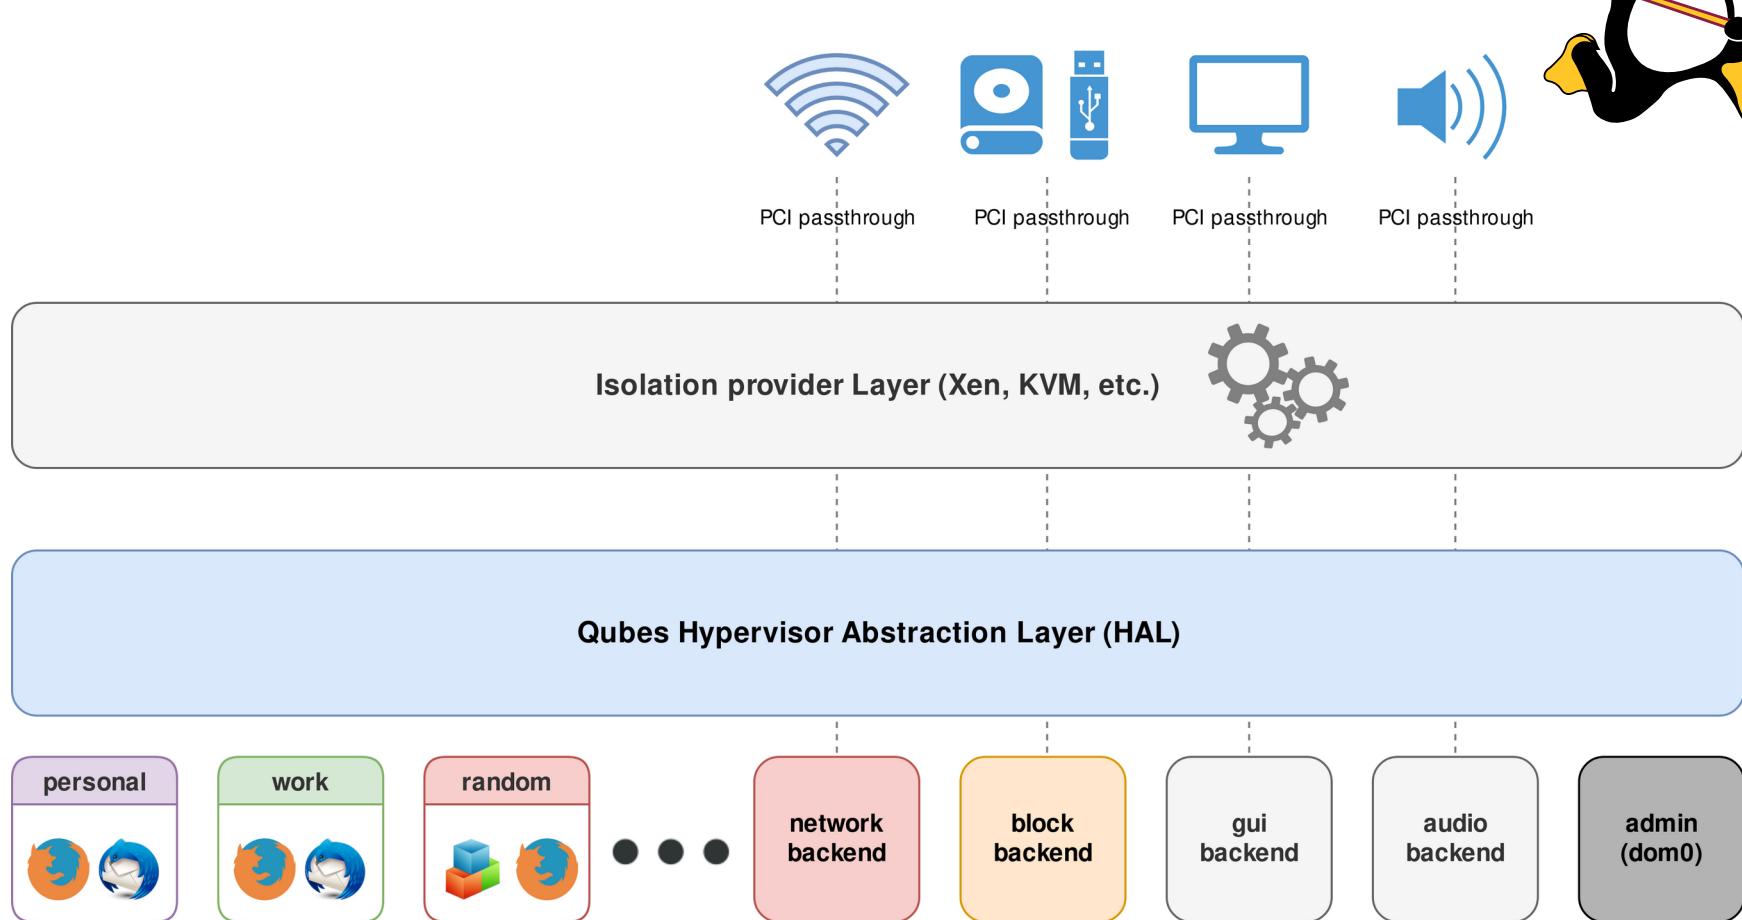
- An operating system based on a hypervisor.
- Everything compartmentalised.
- Separation at every level – Hardware does not interact directly with the end usage (normally).
- All your regular actions are in VMs.
- Entirely open source.

History

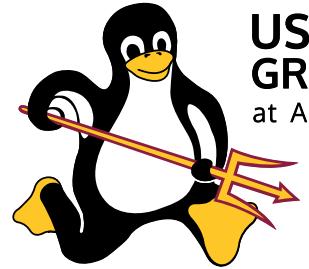
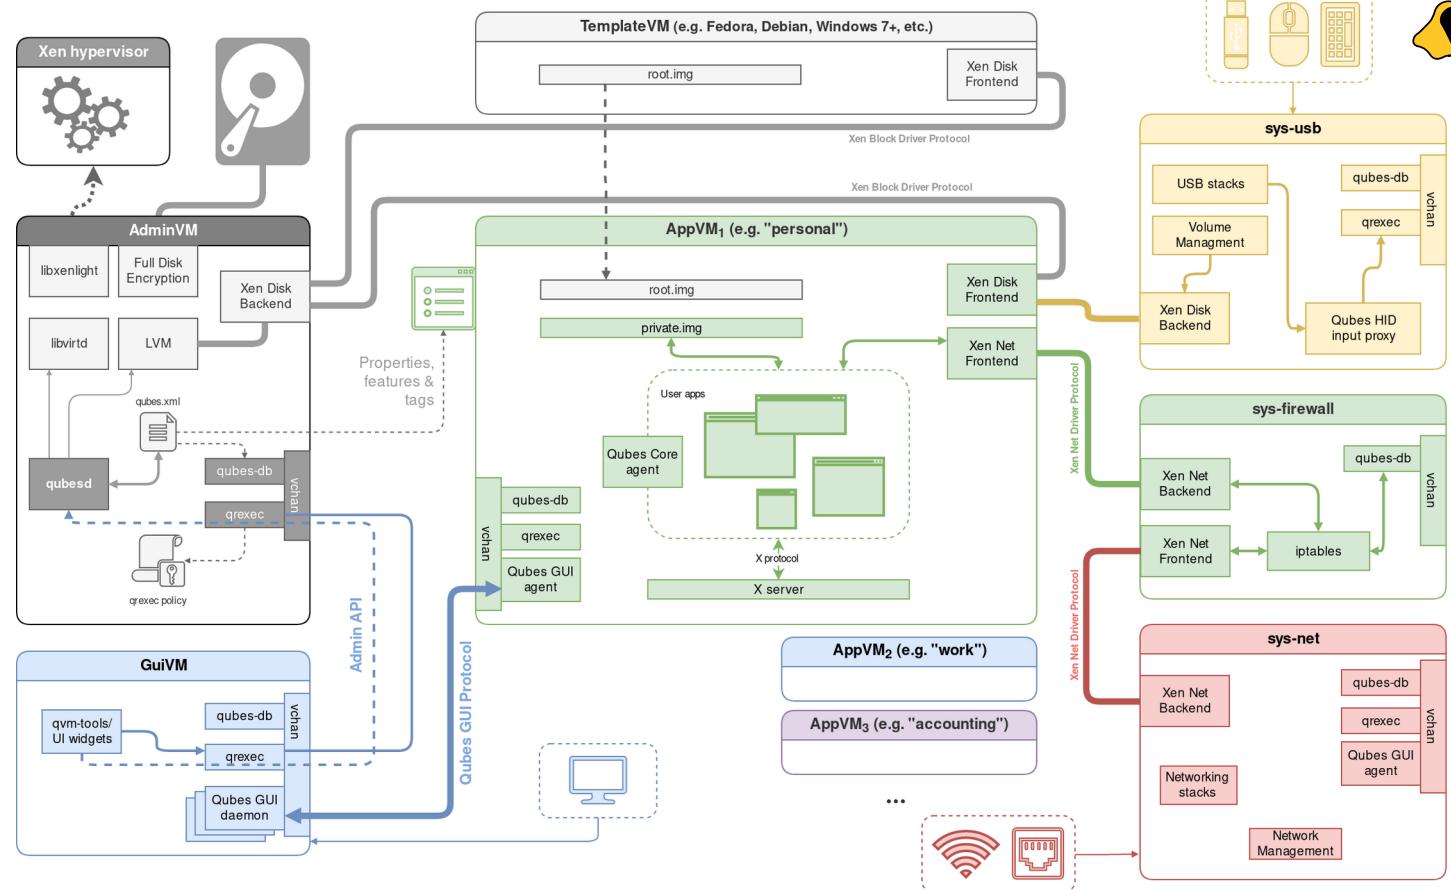
- Introduced in January 2010.
- Founded by Joanna Rutkowska.
- Compartmentalised approach for security.
- Currently led by Marek Marczykowski-Górecki.
- Latest version is 4.1.0 released 4th of February 2021.



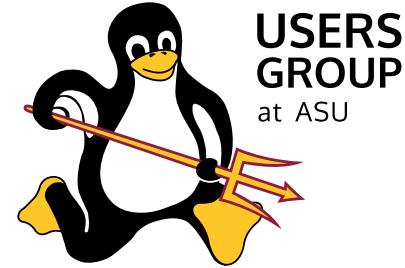
Architecture



Core Stack

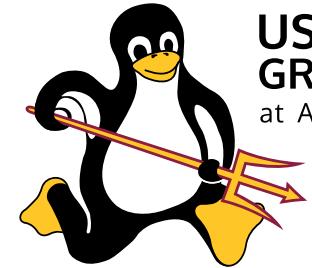
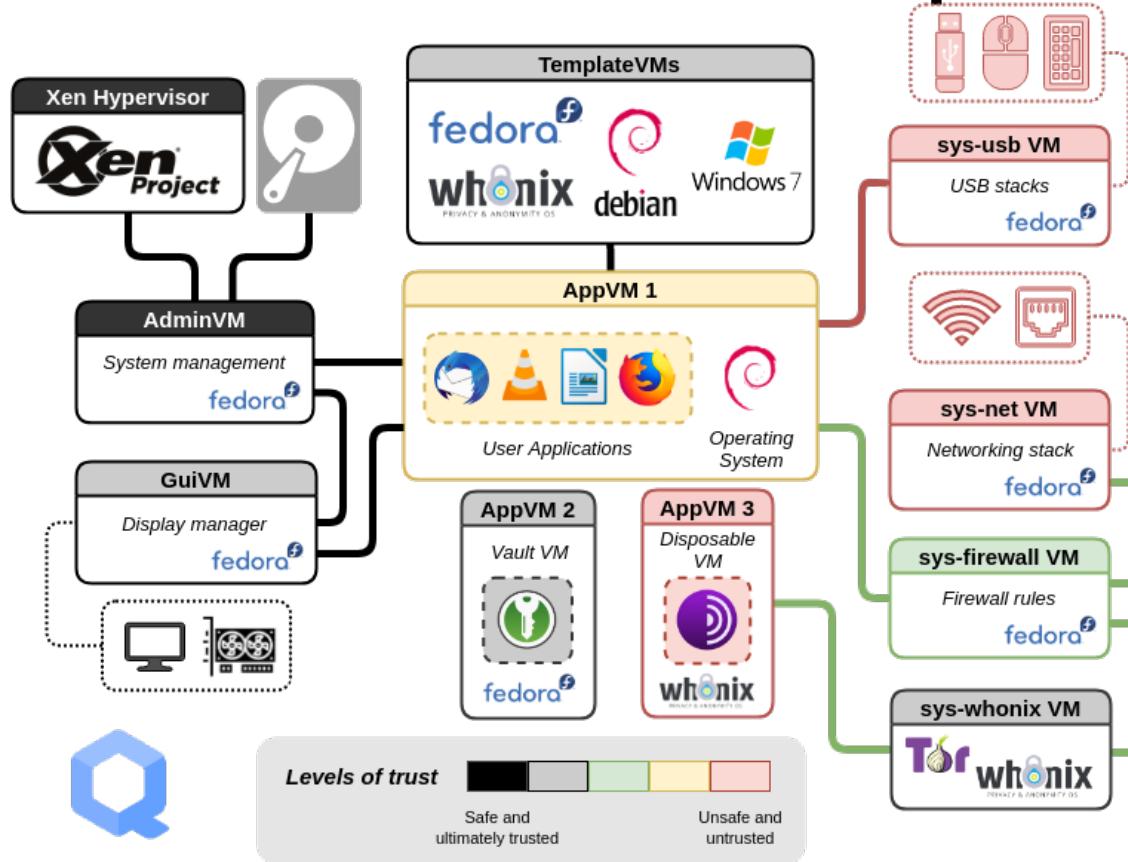


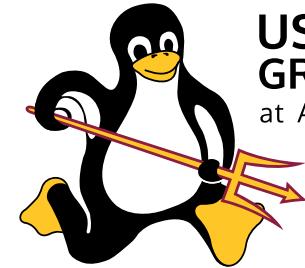
Working in AppVMs



- AppVMs are where you do all your work.
- They can run everything – even Windows
- Connected to the hardware through other VMs.
- Interact with the core and other VMs with drivers.

Use Case Example





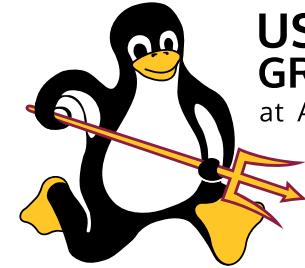
Why?

- Security by Isolation Approach
 - Virtualisation
 - Compartementalisation.
 - Need to Have approach.
 - Disposable VMs.
- Privacy
 - Whonix and Tor
 - Preventing Fingerprinting.
- Compartementalisation of your Workspace.
- Changing Systems Often

Why Not?

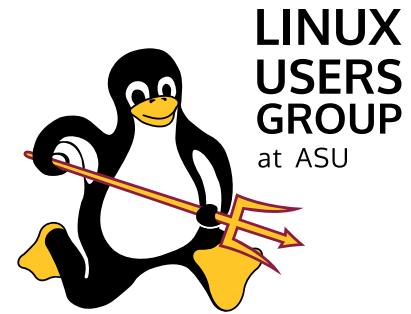


- Technical experience occasionally required.
- Resources
 - Overhead
 - Sharing
 - Device Specifications
- Thinking in a Qubes way
- Technical Restrictions

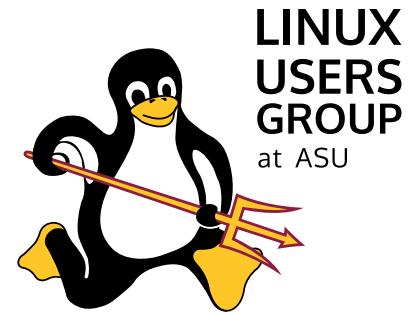


Why do I?

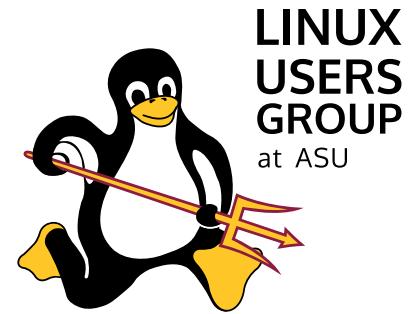
- Compartmentalisation
- Security
 - Isolation
 - Need to Have
- Mirroring across devices
- Whonix



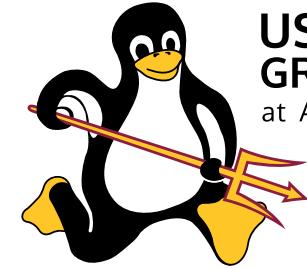
Section 2 – Usage and Security Implementations



Section 2.1 – Virtual Machines in Qubes



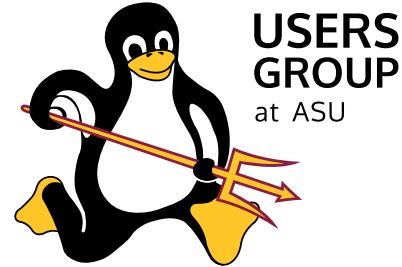
Section 2.1.1 – Virtual Machine Types



Types of VMs

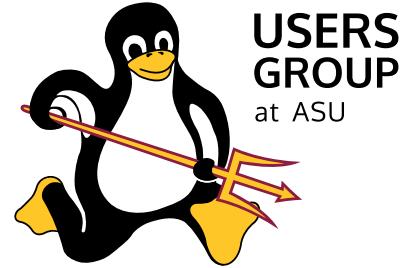
- DispVMs – Based off template and no persistance.
- AppVMs – Based off a template and limited persistance.
- TemplateVMs – Root file system for AppVMs.
- StandaloneVMs – ‘Traditional’ VM.
- AdminVM – The core, dom0.

DispVMs



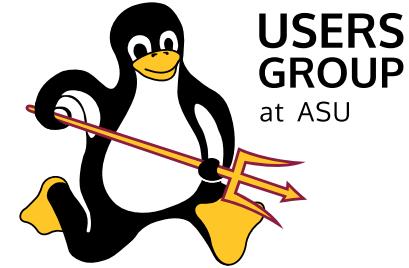
- Disposable VMs are temporary disposable systems.
- Based off a template.
- Inherits the entire template including `/home/`.
- No persistence whatsoever, fully wiped on exit.
- Has a temporary volume though can be set to only use RAM.
- Files can be easily opened to be viewed.
- Randomly numbered between [1-10000).
- **Demo!**

AppVMs

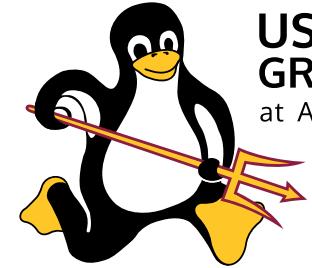


- Based off a template.
- Entire `/rw/` directory persists which includes:
 - `/home/`
 - `/usr/local/`
- Inherits part of its `/home` and `/usr/local`.
- Is provided an image of the template VM which was **generated at its last shutdown**.
- Where you will do all your work.
- **Demo!**

Template VMs

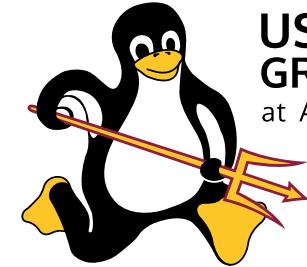


- Templates have no inheritance, recursive templates is not a possibility.
- Everything persists so **be careful**.
- No network by default.
- Its image is updated at shutdown.



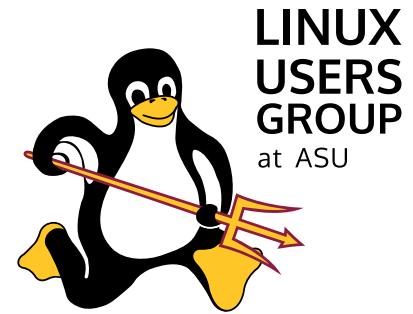
Standalone VMs

- Traditional VMs.
- Either independent clone of template or separately installed system.
- No integration though possible with drivers.



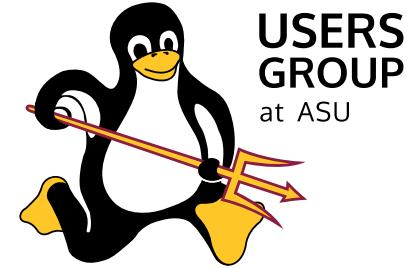
AdminVM

- AdminVM is dom0.
- Is the backbone of the operating system.
- Where the entire backend of the OS is ran.
- No network connectivity.
- Updates can only be made to specific repositories.
- Runs Fedora 32.
- Isolated from the rest of the system – No clipboard integration.

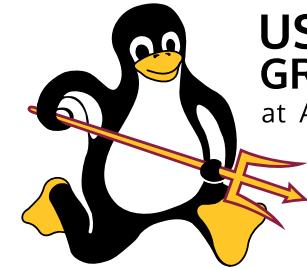


Section 2.2 – Template VMs

Template Repositories



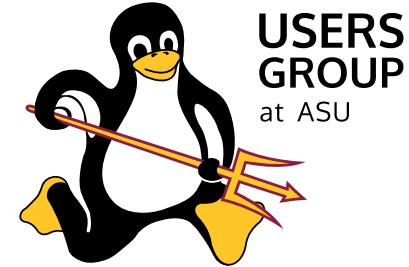
- Official Templates – Built and released by the Qubes team.
- Community Templates – Designed, built and maintained by the community.



Official Templates

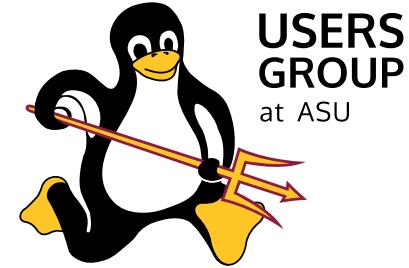
- Fedora 34
 - Minimal Template
 - Xfce Template using Xfce instead of GNOME.
- Debian 11
 - Minimal Template
- Whonix 16 Gateway and Workstation

Community Templates

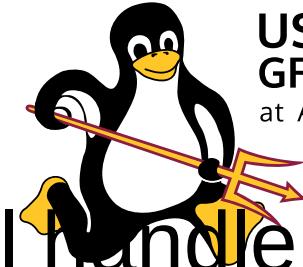


- Whonix
- Ubuntu
- Arch Linux
- CentOS
 - Minimal Template
- Gentoo
 - Minimal Template
- NetBSD

Custom Templates

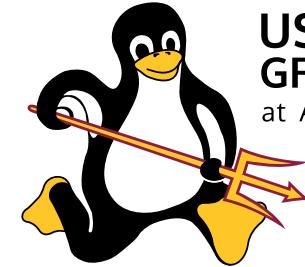


- Templates which aren't standalone but can be implemented.
- Kali Linux
- Parrot



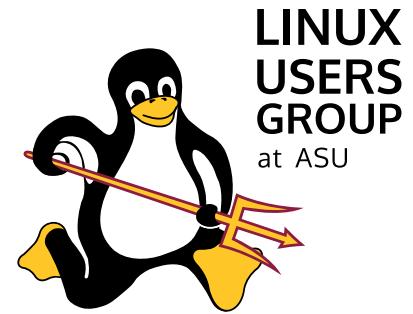
Updating

- Qubes has an updating utility which will handle proper updating for Templates.
 - Creates a mirror.
 - Updates.
 - If successful → Updates the actual Template
- Manually in the individual TemplateVM.

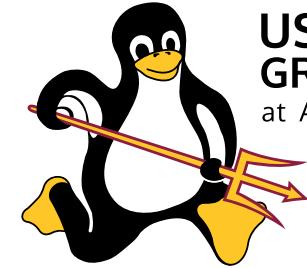


Installing Programs

- Using the package managers.
- Sending over binaries and repositories and installing them locally.
- Adding them manually in directories such as `/opt/` and creating respective desktop files.

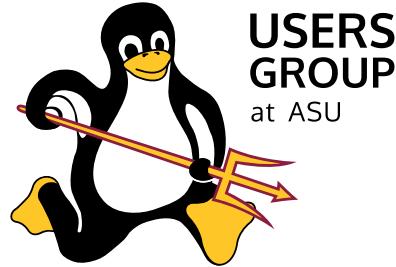


Section 2.3 – Migration and Backups



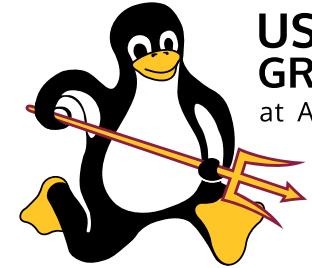
Cloning

- Any VM, regardless of type can be easily cloned.
- Cloning a VM is accessible directly in the basic settings.



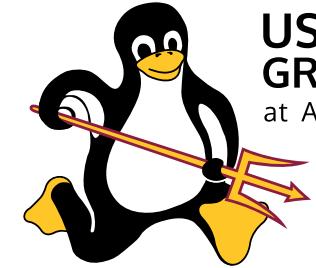
Qubes Backups

- Qubes official backup utility.
- Packages a set of VMs into a backup.
- Saves the version of the last shutdown.
- Saves specific file systems:
 - AppVMs – private.img
 - TemplateVMs – root.img
 - StandaloneVMs – root.img and private.img
- Encrypts backups with scrypt.



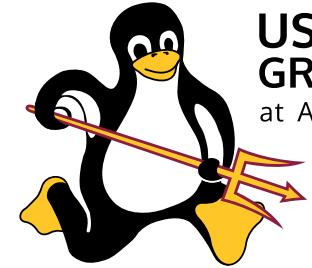
Backup Encryption

- Uses the scrypt encryption utility.
- Initially developed for the Tarsnap online backup system.
- Uses a complex cryptographic system including both signing and encryption, both which are required for decryption.
- Qubes uses their release signing key for the signing.



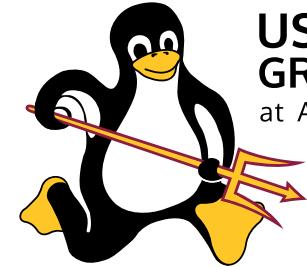
Tarsnap Cryptography

- A 2048-bit RSA key used for signing archives. This is used in combination with SHA256 and a Merkle hash tree to verify the authenticity of stored archives.
 - Cryptographers: RSASSA-PSS is used, with SHA256 as the hash function.
- A 2048-bit RSA key used for encrypting session keys. All the data which the Tarsnap client sends to the server to store is encrypted with per-archive random AES-256 keys; those keys are encrypted with this RSA key and attached to the stored data.
 - Cryptographers: RSAES-OAEP is used to encrypt the session keys, using MGF1 and SHA256, and the padding verification performed when decrypting is carefully written to be free of timing side channels. The AES-256 keys are used in CTR mode, with sequentially incrementing nonces — generating a new AES-256 key for each session ensures that a key-nonce pair will never be used twice.
- A 256-bit HMAC-SHA256 key used to protect each individual block of data from tampering. From a cryptographic perspective, this is unnecessary, since a Merkle hash tree protects each archive; but data is compressed using zlib before being stored, so this provides protection against a theoretical attacker who can tamper with stored data and has found a security flaw in zlib decoding.
 - The zlib decoding process is quite complex — for that matter, the same is true of all decompression algorithms — and such complex code has a relatively large risk of having security vulnerabilities, hence the decision to add an extra layer of security here.
- Two 256-bit HMAC-SHA256 keys used to generate names for blocks of data stored. Tarsnap uses the same reference-by-hash trick as the author's Portsnap and FreeBSD Update utilities; using HMACs instead of raw SHA256 hashes prevents any information from leaking via the hashes.
 - Why two keys? One is used to hash data, and the other is used to hash archive names. Yes, that's right, even the names of archives are stored securely!



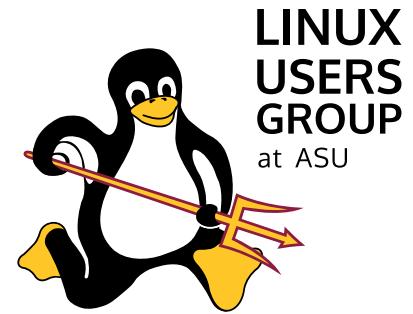
Backup Restore

- Official Qubes Restore utility.
- Restores backups without a hitch.
- Has the possibility to solely verify integrity.
- Can either restore alongside or overwrite.
- Restoration has no difference machine to machine as long as they are on the same Qubes release.

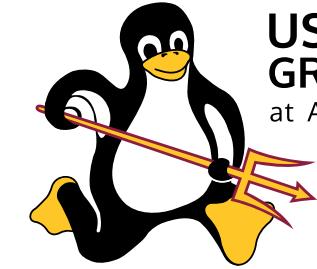


Migration to an ISO

- Official support with Qubes Builder
- Allows you to turn a Qubes image into an ISO
- Theoretically should be flashable onto another computer just like any ISO.

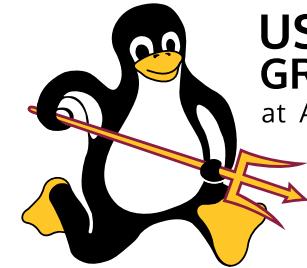


Section 2.4 – Interactions between VMs



Clipboard Interactions

- A single elevated clipboard.
- Any VM can send their clipboard to the elevated clipboard with `Ctrl+Shift+c`.
- Any VM can retrieve the contents of the elevated clipboard with `Ctrl+Shift+v`.
- Elevated clipboard is wiped on retrieval.
- AdminVM is not integrated with the elevated clipboard.

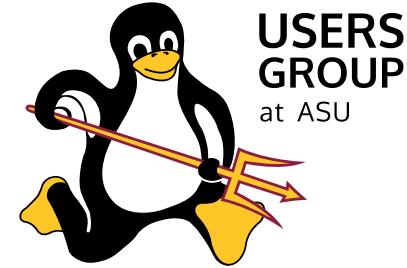


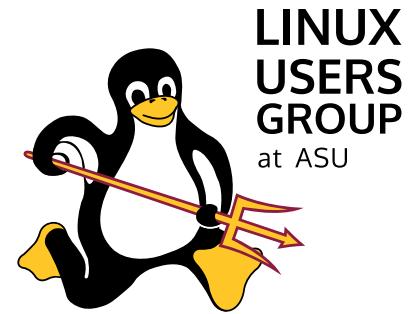
File Interactions

- All VMs can send files to another VM.
- All VMs except the AdminVM can receive files.
- Files will arrive in the `~/QubesIncoming/VM_NAME/` folder.
- AdminVM can only receive files from other VMs by passing the IO from the VM to the AdminVM.
 - i.e. `qvm-run --pass-io QUBE_NAME 'cat NAME_OF_FILE' > NAME_IN_DOM0`

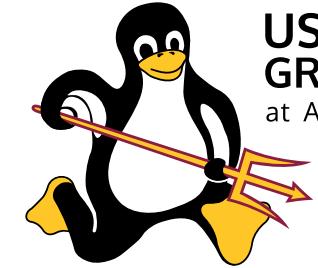
Split Utilities

- Split GPG
- Split SSH
- Split Electrum Wallet
- Split dm-crypt
- U2F Proxy



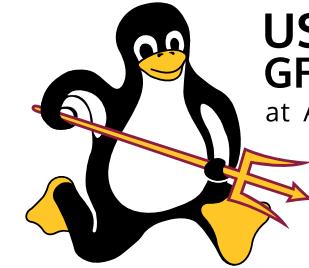


Section 2.5 – Networking



Networking Overview

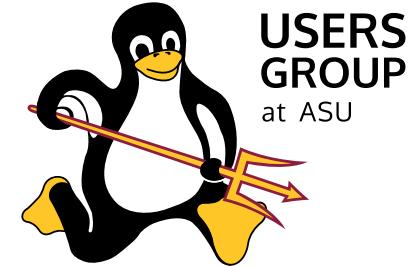
- Any VM can be set to be a networking VM.
- There are two primary VMs responsible:
 - Hardware VM (i.e. WiFi, Ethernet, etc.)
 - Firewall VM
- Other VMs can be configured to handle and adapt a network connection.
- Network VMs can be chained.
- Changing networks is as simple as changing one setting and is applied instantly.
- You select the final VM as the one you want to use.



Hardware VMs

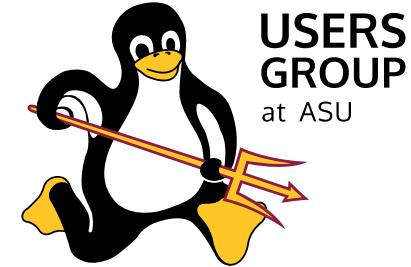
- Hardware networking VMs are those that interact with the actual PCI devices.
- Good practice is one per PCI device.
- Lets everything through.
- Uses NetworkManager by default.

Firewall VM

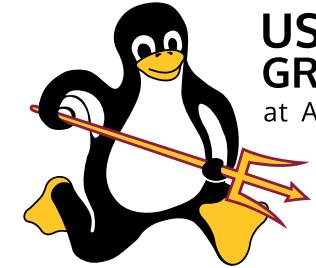


- Firewall rules can be set for each Qube.
- Applies all network-level policies.
- Responsible for blocking inter-VM networking.

Custom VMs

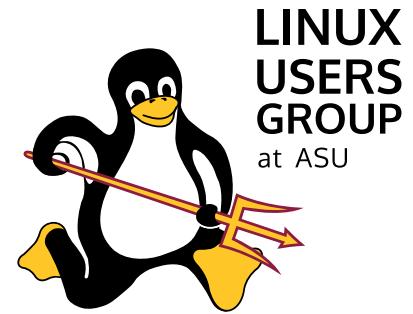


- No limit in the amount of network VMs that can be created and made.
- You can create a VPN VM and everything that is passed through goes through the VPN.
- You can deploy a full on DNS sinkhole (i.e. Pi-hole) or a firewall if you wanted and chain it.

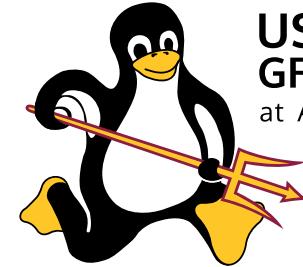


Whonix Networking

- sys-whonix VM running Whonix Gateway.
- All connection passed through it are passed through Tor.
- Tor connection is fully configurable.
- Just like any other VM.

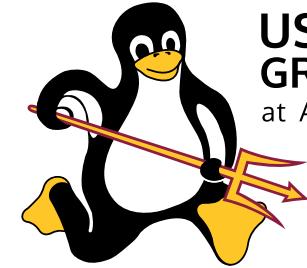


Section 2.6 – USB Handling



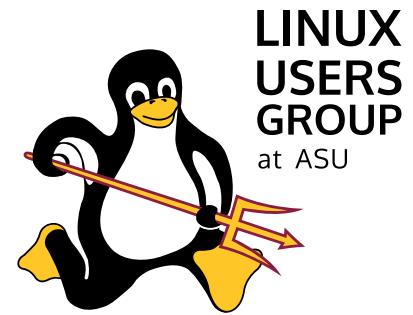
USB Qubes

- Separate VM which interacts directly with the USB PCI device.
- Without a USB Qube, the USB PCI device interacts directly with dom0.
- To use or have access to a USB device in a VM, it HAS to be manually passed through the Qubes Devices taskbar menu.
- **Demo!**



Complications

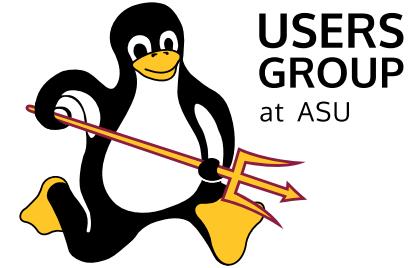
- Necessary passthrough for login
 - Passthrough of USB keyboard device to dom0 on boot through a policy.
- Connection not as stable as if directly passed through.



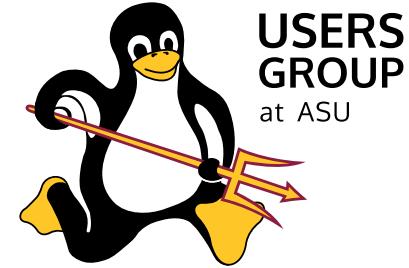
Section 2.7 – VM Configuration

Basic Settings

- Name and label colour
- Template
- Networking VM
- Private storage size
- Start on boot
- Include in backups
- Debug mode



Advanced Settings



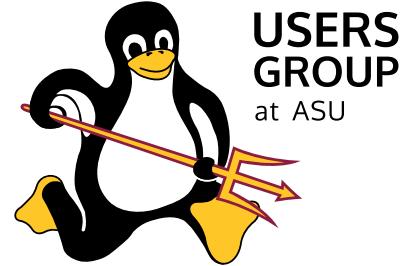
- Initial Memory
- Max Memory (only if memory balancing is enabled)
- VCPUs Number
- Kernel
- Virtualisation
- Disposable VM Template (if not itself)
- Provides Network Toggle
- Boot from ISO

Firewall Rules



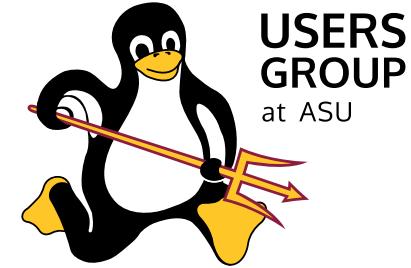
- Configure firewall rules.
- Configurable by:
 - Address
 - Port
 - Protocol
- Temporary Access

Devices



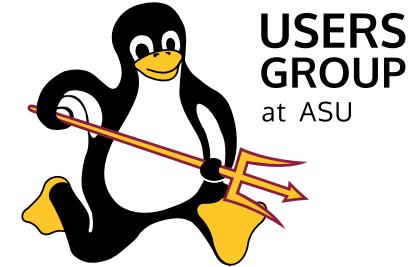
- PCI Passthrough
- Must use HVM or PV virtualisation and have disabled memory balancing for ideal results.
- Limit in types of PCI device that can be passed depending on the VM type.

Applications

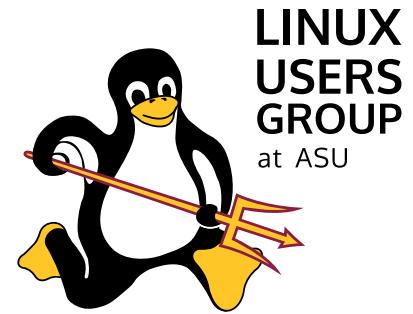


- Select which applications show up in the application menu for that VM.
- Grabbed from all `*.desktop` files in the `/usr/share/applications/` directory.
- Refreshed from Template VM.

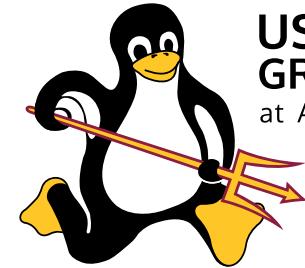
Services



- Add or remove specific services.
- Enable or disable specific added services.



Section 2.8 – Customisation



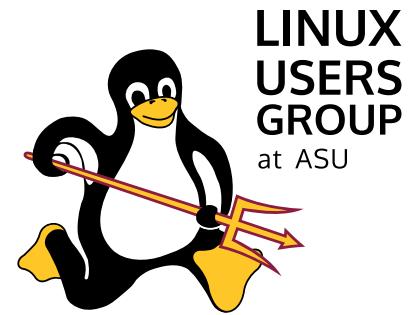
DEs and WMs

- Desktop Environments officially supported:
 - Xfce (default)
 - KDE
- Window Managers (community support):
 - i3
 - AwesomeWM
 - bspwm (unofficial)

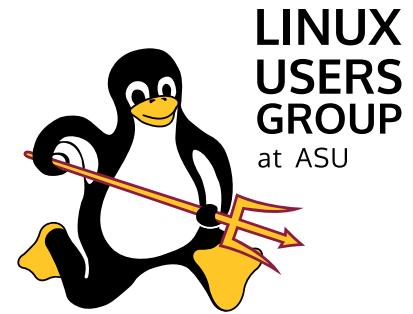
Security when Ricing



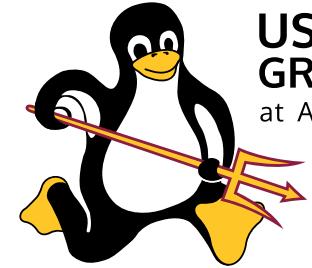
- Passing programs to dom0
- Affecting dom0
- Restriction in availability of programs.



Section 3 – Drivers, Interfaces and Policies



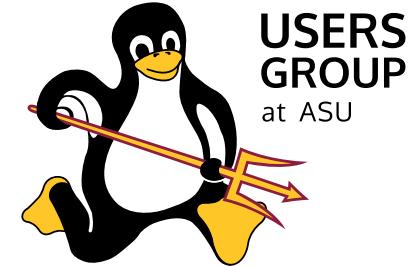
Section 3.1 – Intercommunication Drivers



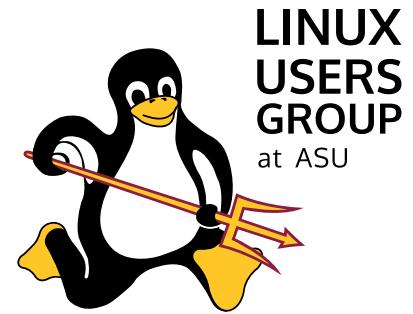
Clipboard

- A single elevated clipboard.
- Any VM except dom0 can copy and paste from it.
- It can only store a maximum of one entry.
- When it is pasted in a VM, it is wiped.

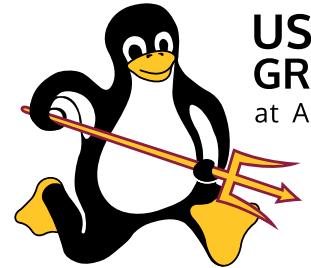
File Transfer



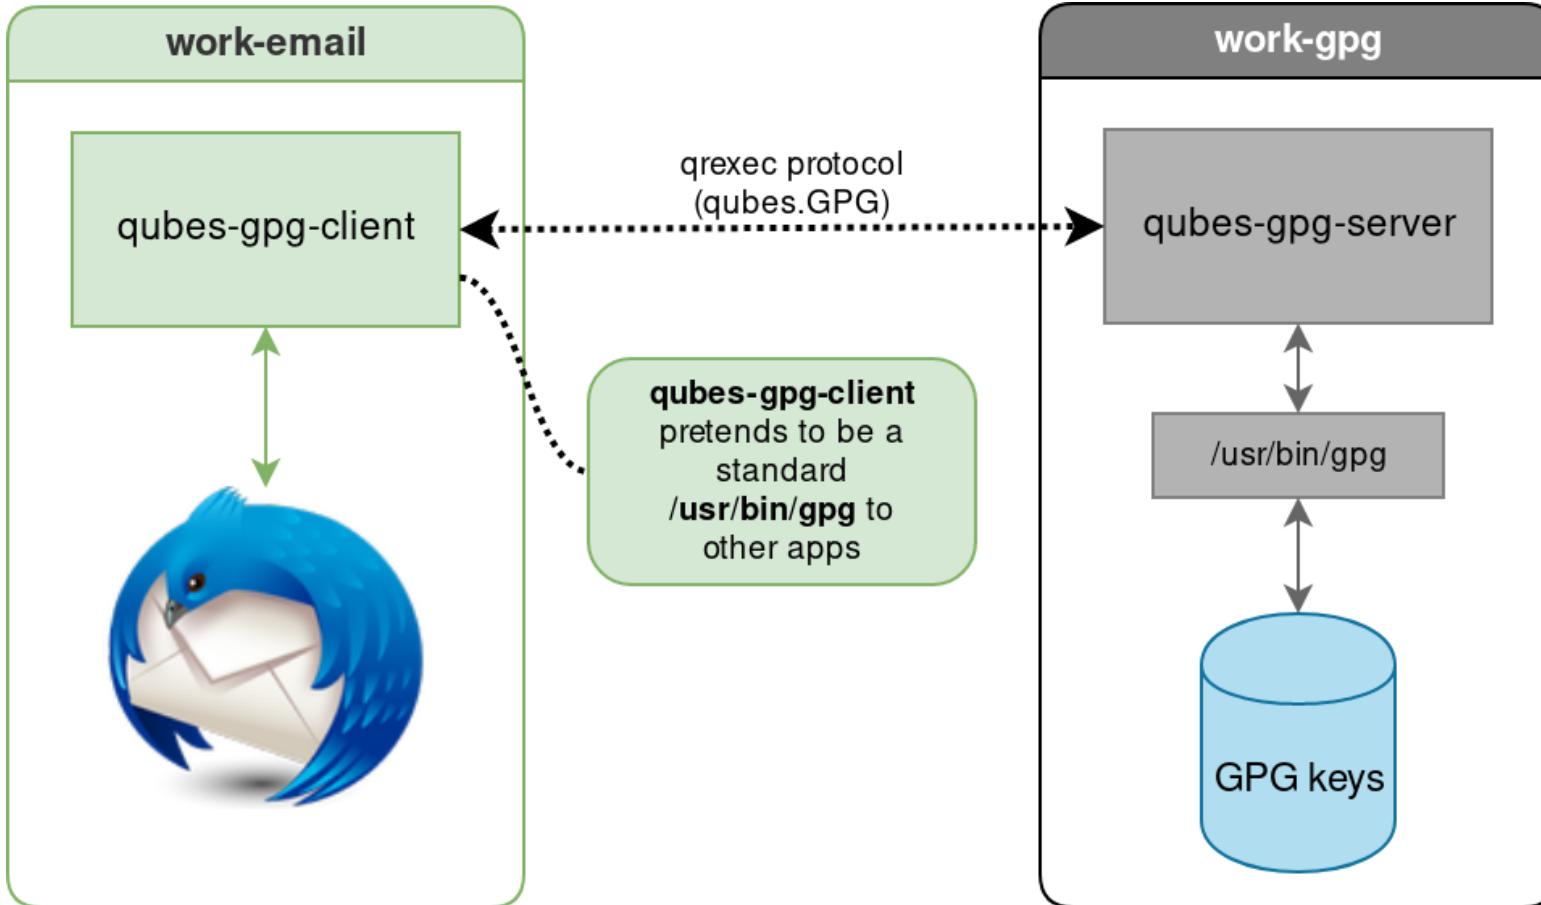
- Files can be transferred between all VMs except dom0.
- Files will be sent to the `~/QubesIncoming/VM_NAME/` directory.

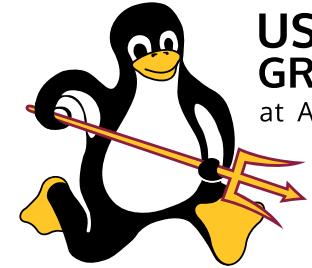


Section 3.2 – Split Interfaces



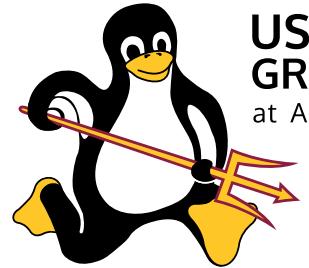
Split GPG - 1





Split GPG - 2

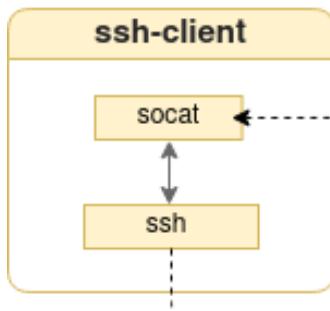
- All keys are stored in a separate VM which is ideally fully isolated.
- The qubes-gpg-client is what interacts with the client, passing all the commands over to the GPG program on the GPG Qube and sending its output back.
- Two sets of prompts appear at each GPG usage.
- Access can be determined to a set amount of seconds.
- Passwords are entered and validated directly in the GPG Qube.
- You set your GPG client to be the Qubes client in programs.
- **Demo!**



Split SSH - 1

Plain Setup

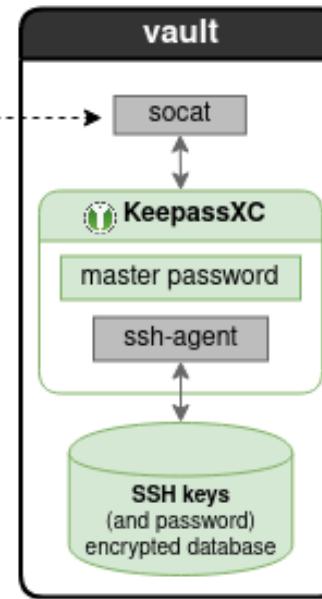
uses ssh-askpass



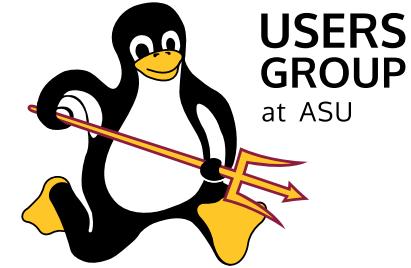
Remote Server

KeePassXC Setup

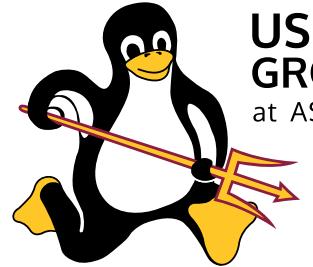
Uses the password manager to protect your keys



Split SSH - 2



- SSH keys safe in an isolated VM.
- SSH keys are retrieved and passed through when requested.
- SSH private keys are kept in the isolated VM.
- Unfortunately, use can be held on for the entire of the existence of the VM.
- Community made.



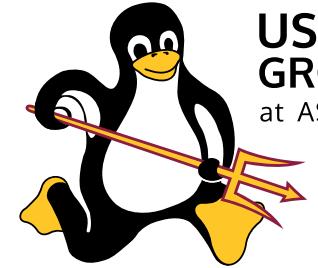
Spit dm-crypt

- Splits up the process of decryption of a LUKS1 partition.
- Encrypted VM is attached to a specifically configured offline disposable VM.
- Prompts for a password and decrypts the partition in the disposable VM.
- Mounts the partition into the destination VM.
- The destination VM does not have access to the password or encryption key.

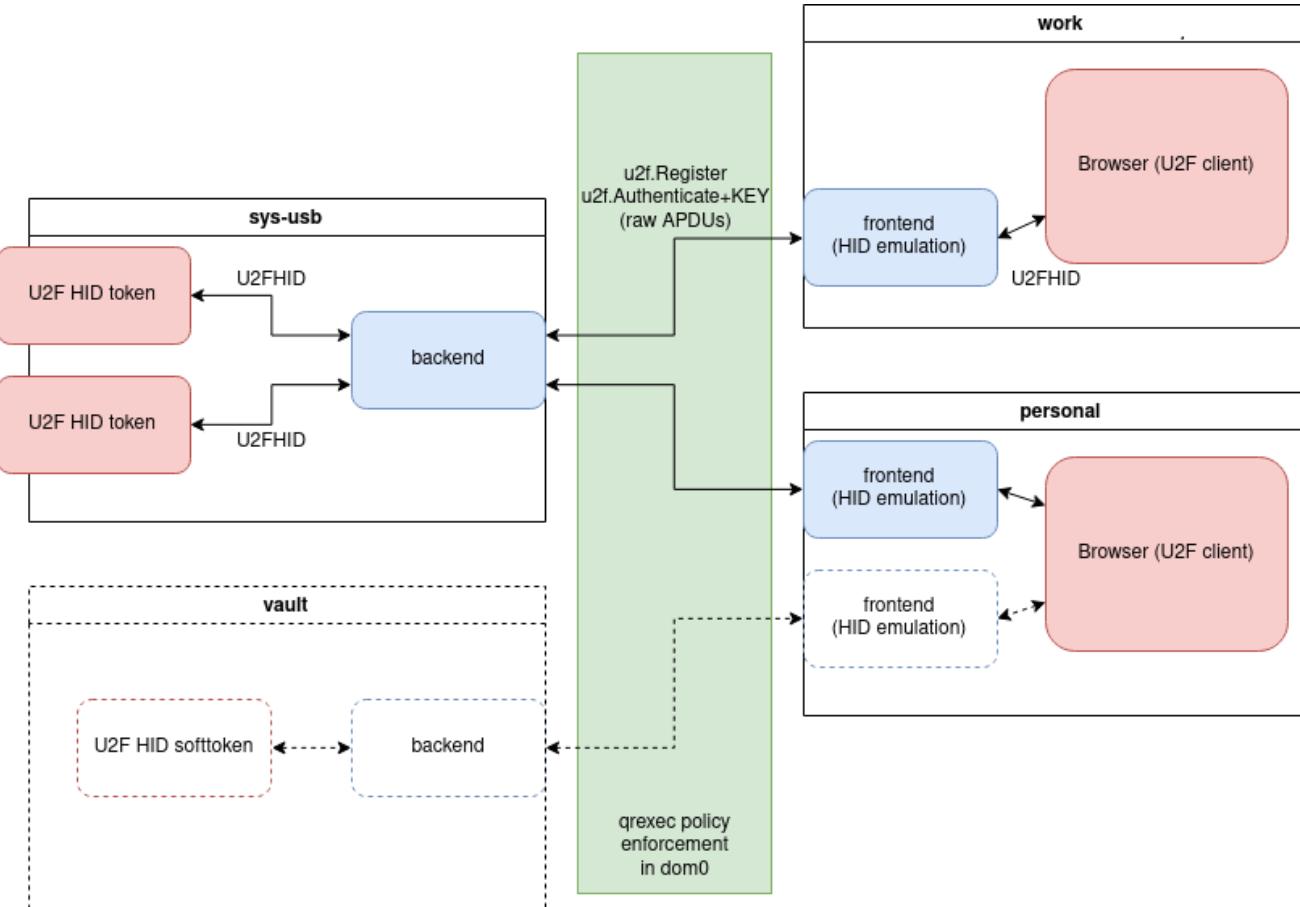
Split Electrum Wallet



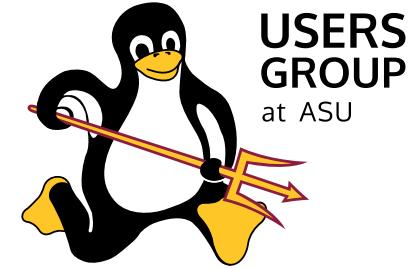
- Wallet is stored in one isolated offline VM.
- Wallet is accessed via an online watching-only Electrum wallet which connects to the isolated wallet.



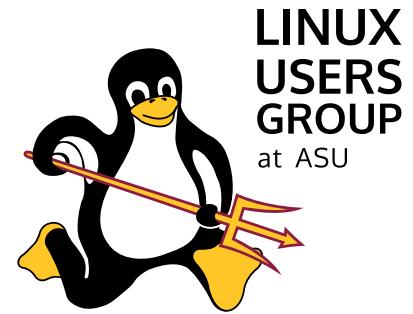
U2F Proxy - 1



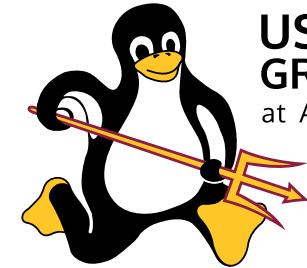
U2F Proxy - 2



- The proxy only passes the necessary to the actual U2F key, performs the transaction and returns it to the browser without passing the entire U2F key to the device.
- Can be set to only allow specific tokens per VM.
- No direct interaction between the U2F device and the VM that is requesting the token.



Section 3.4 – Policies



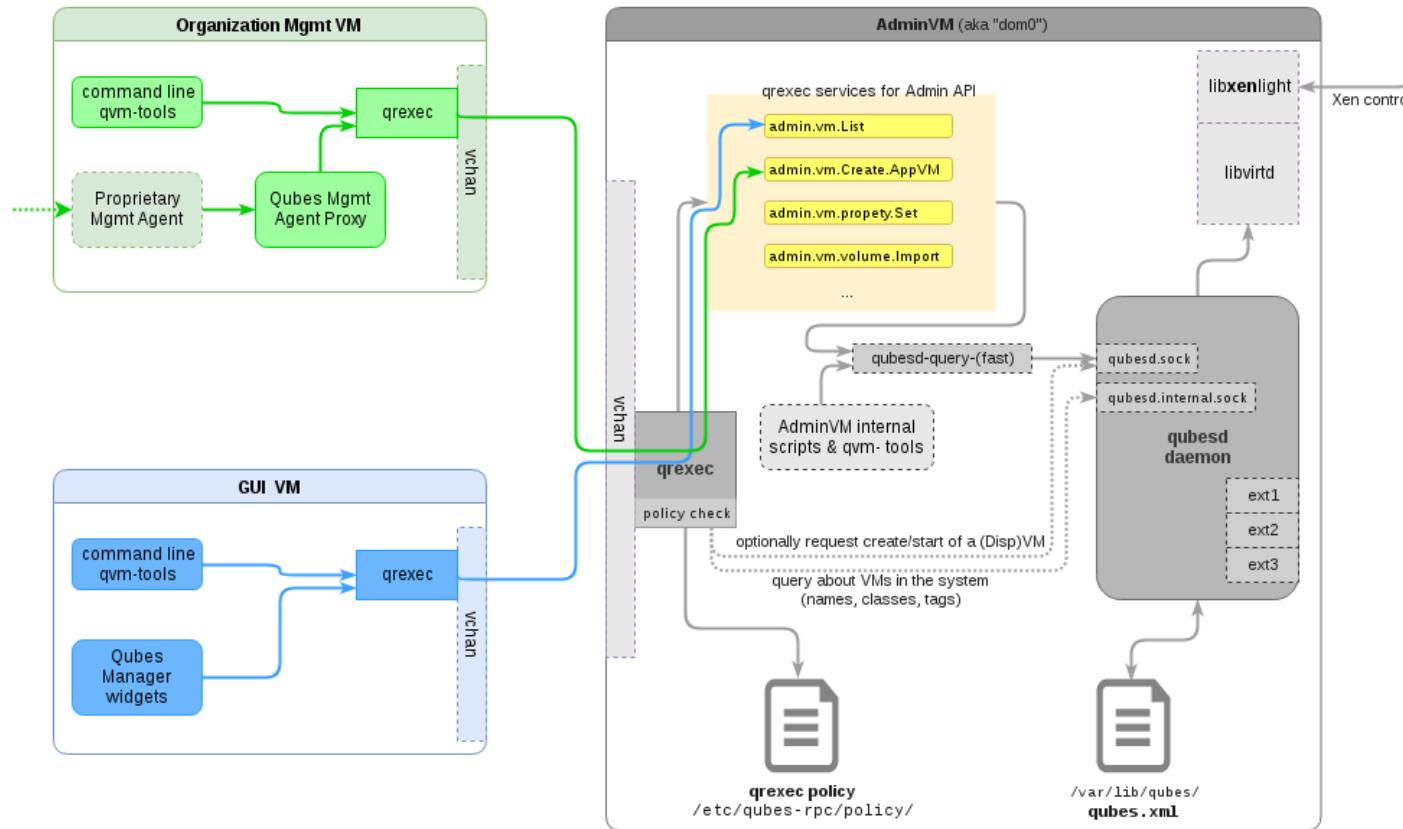
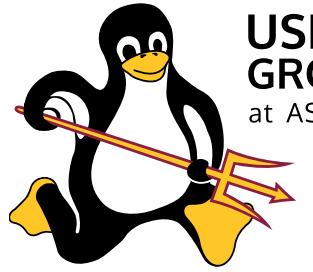
Policies in Qubes

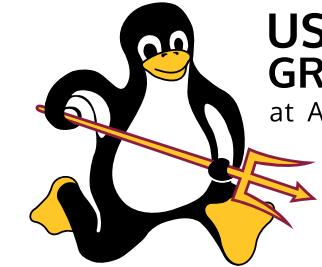
- Each interaction with a driver is controlled by a policy.
- Policies can be set both inside of a VM for its drivers or for the system.
 - `/rw/config/`
 - `/etc/qubes-rpc/policy/` in dom0

Control of Policies

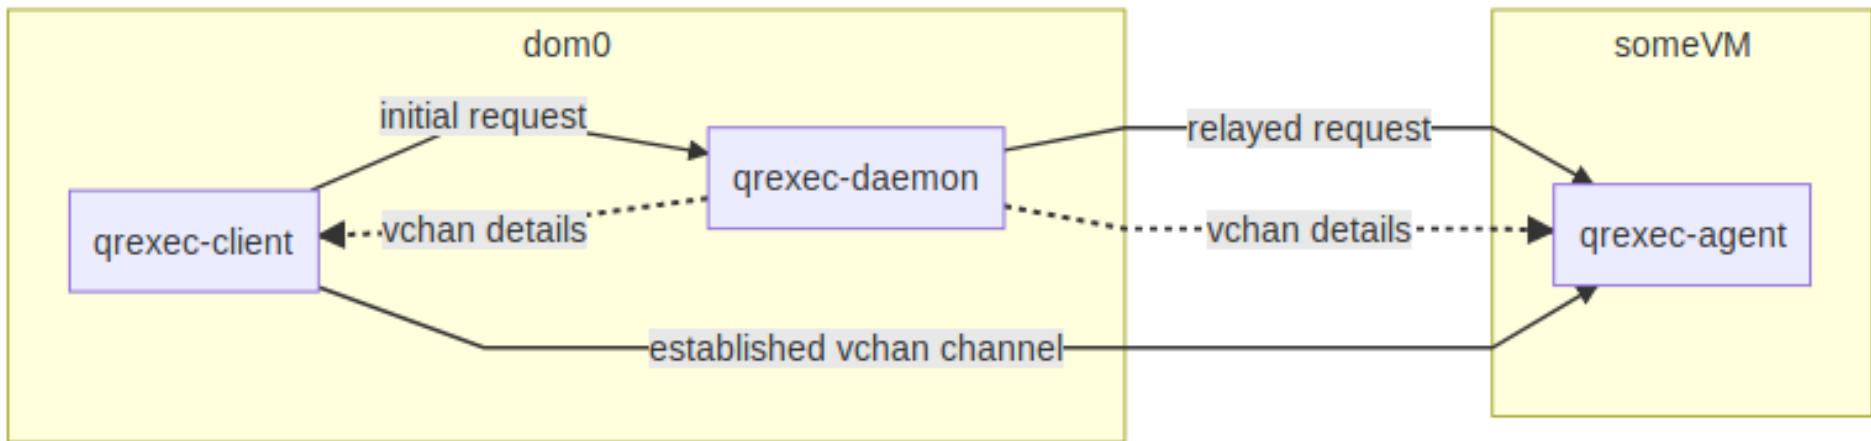


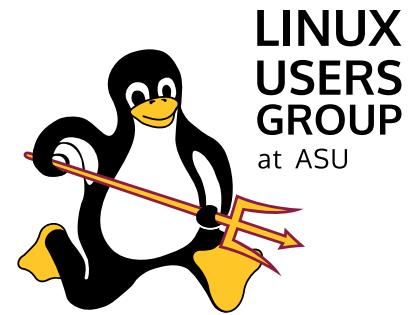
- Which VMs does the policy apply to.
- How to handle any VM not explicitly referred to.
- Handle policies depending on the tag of a VM.
- Allow or deny specific actions.



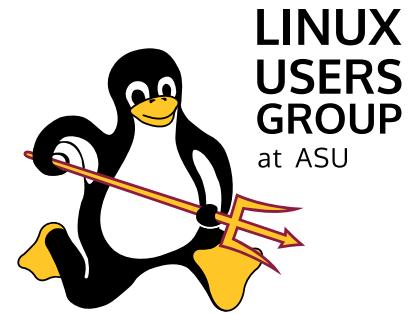


Policy Client Overview

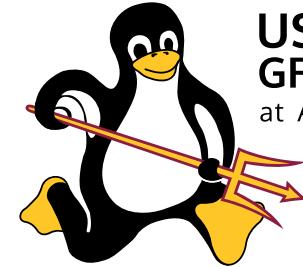




Section 4 – Technical Analysis of the Architecture



Section 4.1 – Volume System



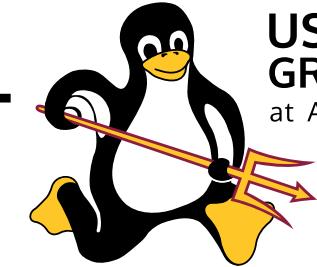
Volume Types

- root.img – Root device (`/`).
- private.img – Persistant volume of a VM.
- volatile.img – Swap and temporary `/`. Discarded on reboot.
- modules.img – Kernel modules and firmware.

Block Devices

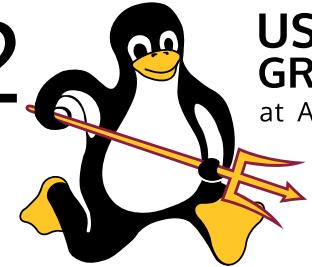


- xvda – Base root (**NOT** root.img)
- xvdb – private.img
- xvdc – volatile.img
- xvdd – modules.img

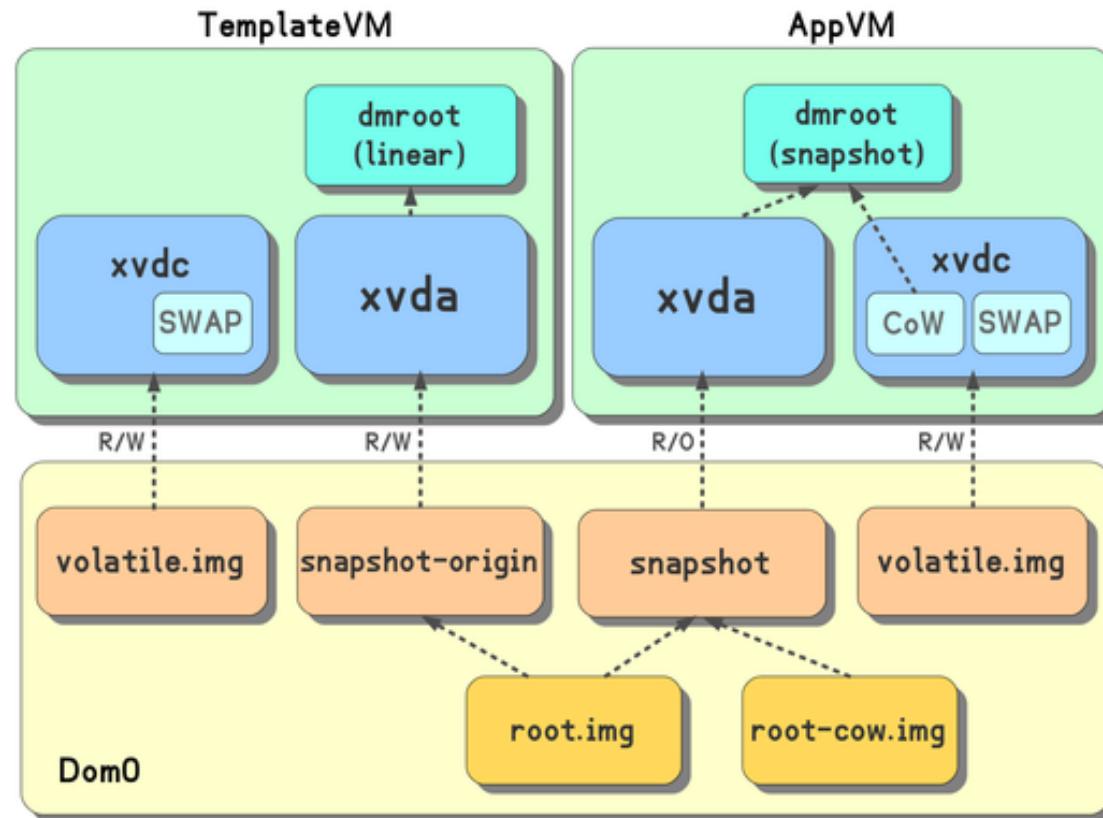


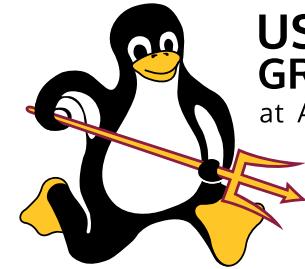
Root Implementation - 1

- Two versions of `root.img`:
 - `root.img` – True root filesystem.
 - `root-cow.img` – Snapshot of the root filesystem.
- Everytime an AppVM is started, a `root-cow.img` is made (unless no changes).
- AppVMs are based off the `root-cow.img`.



Root Implementation - 2



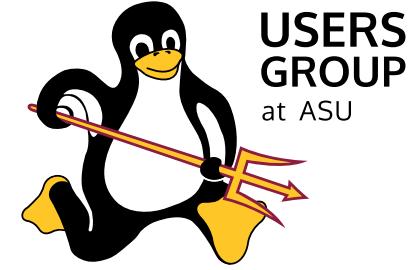


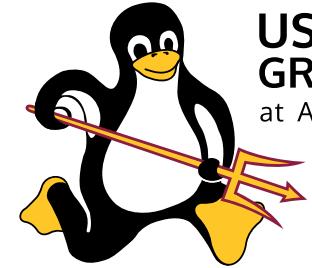
private.img

- Persistent storage for an AppVM.
- What is mounted in `/rw/`.
- Mounted in read-write.
- Contains:
 - `/home/` – Bind mounted on `/rw/home/`
 - `/usr/local/` – Symlinked to `/rw/usrlocal/`
 - Qubes configuration files in `/rw/config/`.

volatile.img

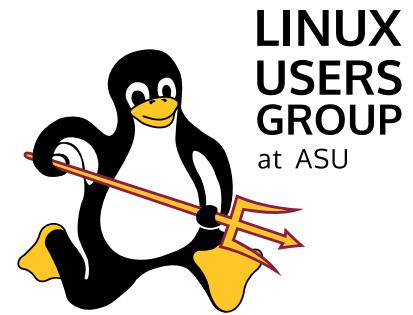
- Contains temporary changes.
- Discarded after AppVM shutdown.
- Contains two primary elements:
 - Changes to the root partition.
 - Swap partition.



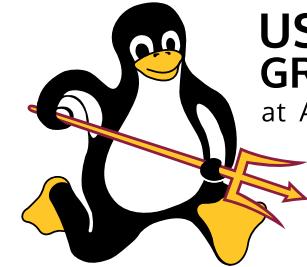


TemporaryVMs

- TemporaryVMs are a unified version.
- Each have their own `root.img` and `private.img` which act in the same way except both are persistent.
- Still use a `volatile.img` for swap.

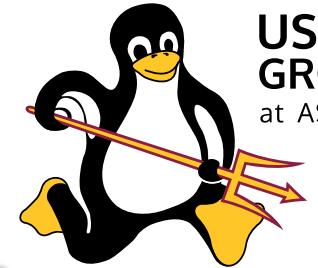


Section 4.2 – Virtualisation

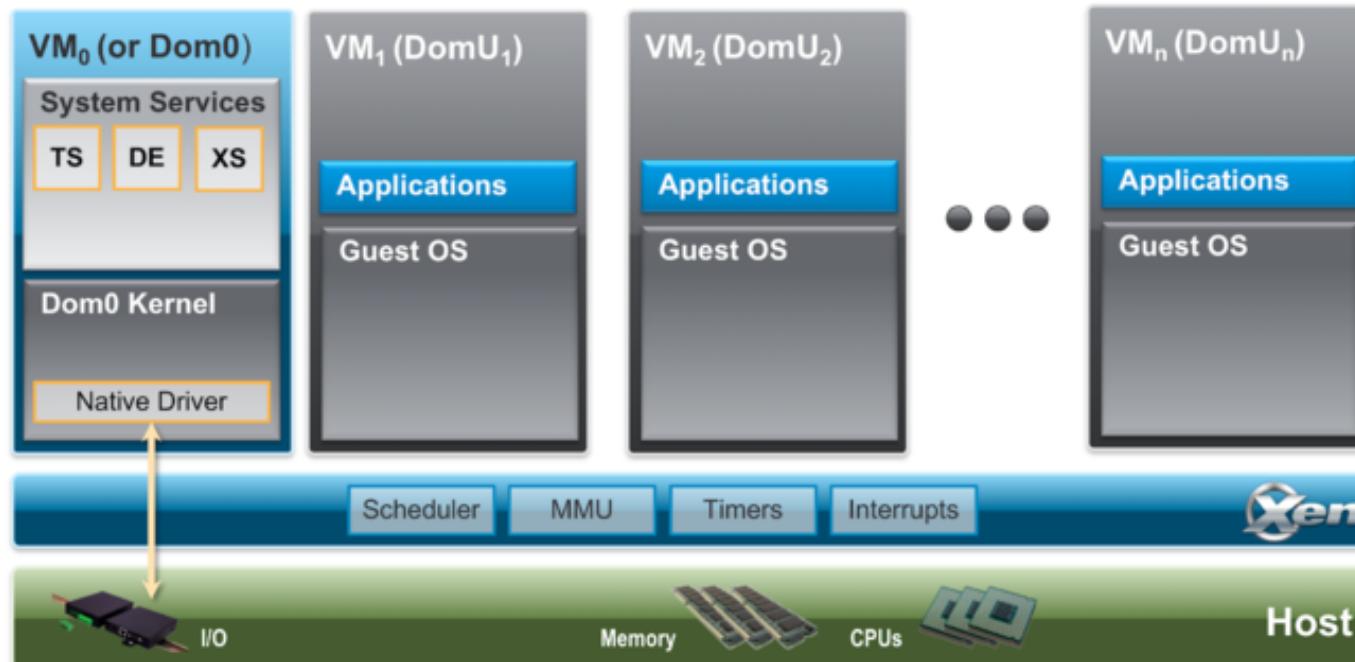


Hypervisor

- Qubes uses the Xen Project Hypervisor.
- Xen is an open source type 1 hypervisor.
- Qubes architecture founded on the Xen Project Architecture



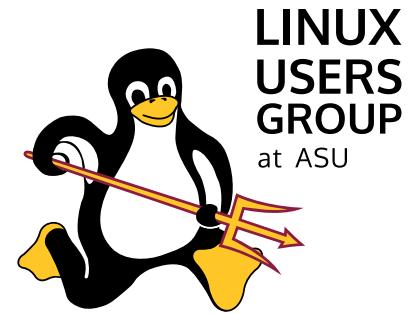
Xen Architecture



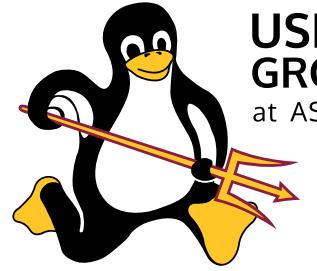
VM Virtualisation



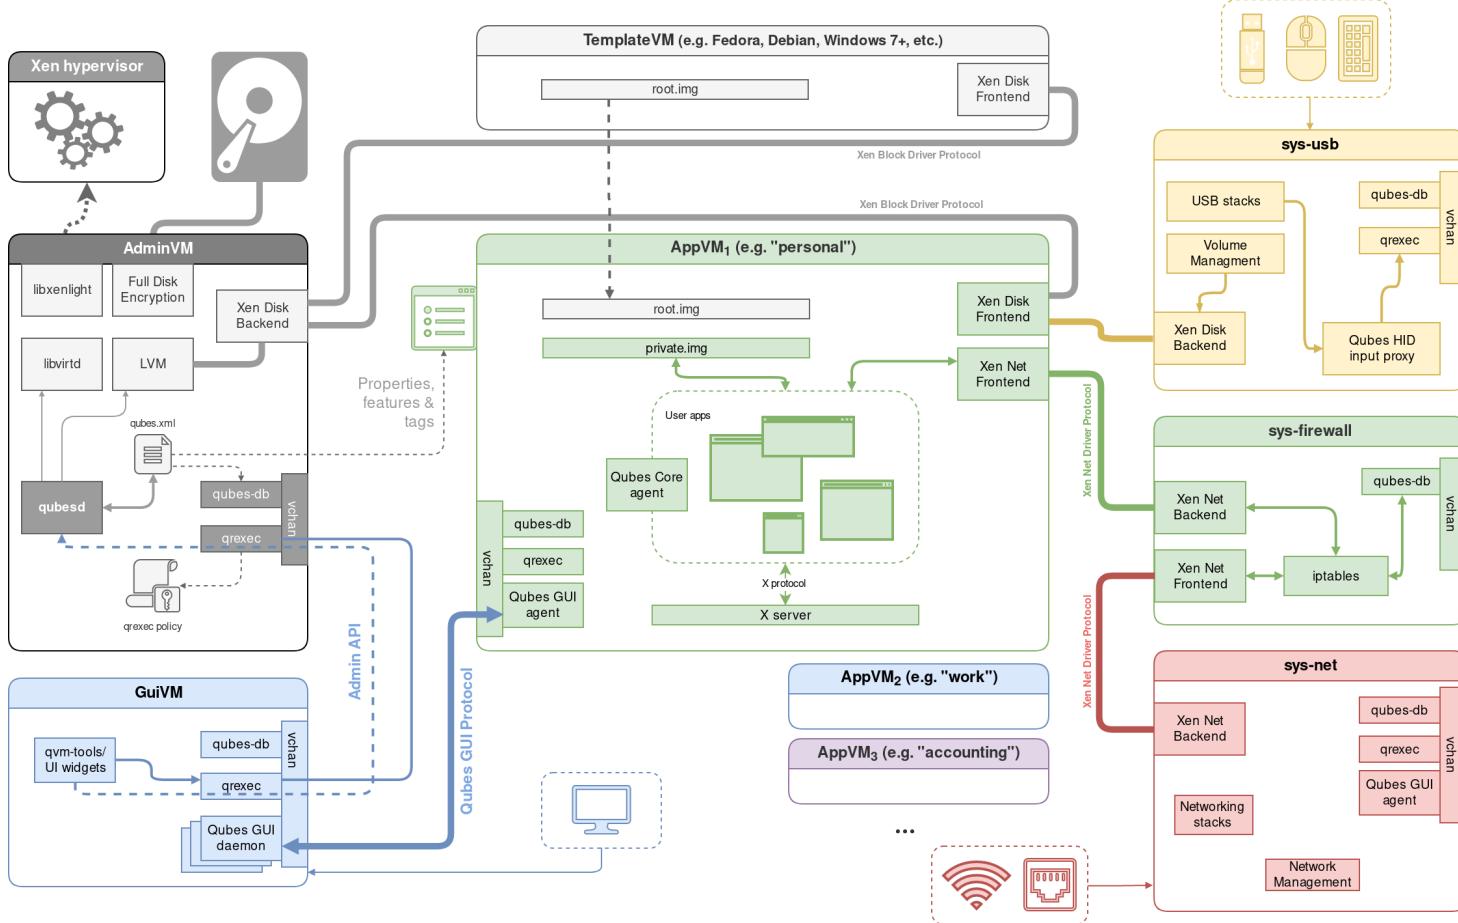
- Virtual Machines can be virtualised in three virtualisation types:
 - PVH (default)
 - HVM
 - PV



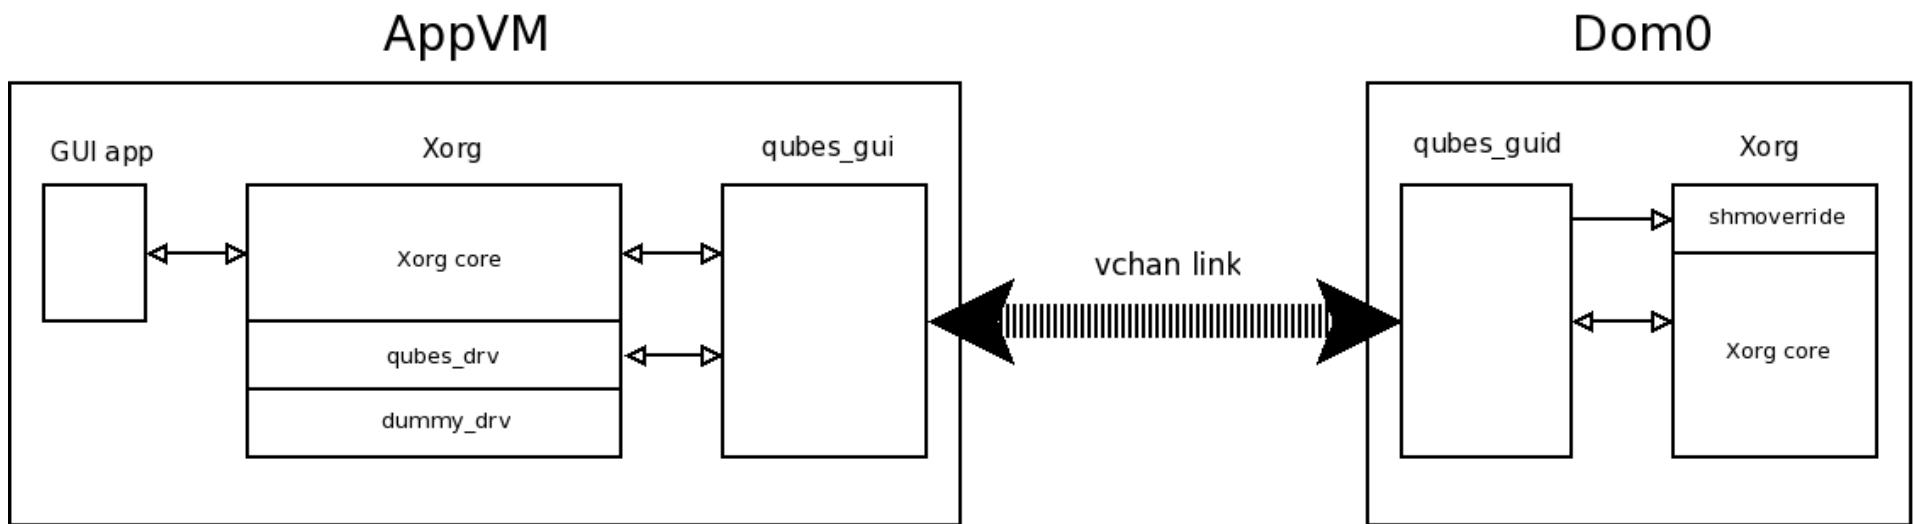
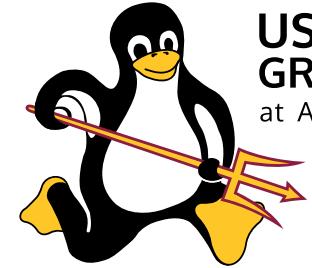
Section 4.3 – GUI



Architecture Overview



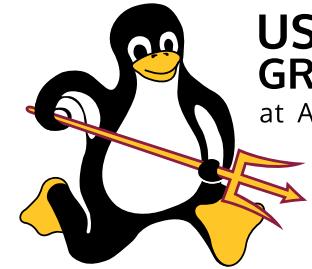
GUI Protocol Overview



GUI Drivers

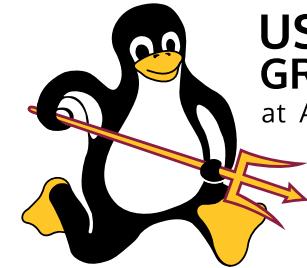


- “Hardware” Drivers
 - dummyqsb_drv – Video driver that paints onto the framebuffer in the RAM.
 - qubes_drv – Provides a virtual keyboard and mouse.
- Qubes Drivers
 - qubes_gui – Runs in the AppVM.
 - qubes_guid – Runs in dom0.



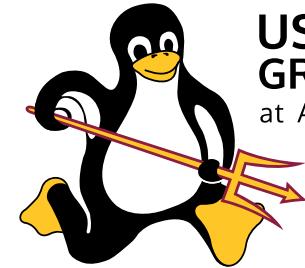
qubes_gui Driver

- Call XCompositeRedirectSubwindows on the root window, so that each window has its own composition buffer.
- Instruct the local Xorg server to notify it about window creation, configuration and damage events; pass information on these events to dom0.
- Receive information about keyboard and mouse events from dom0, tell qubes_drv to fake appropriate events.
- Receive information about window size/position change, apply them to the local window.



qubes_guid Driver

- Create a window in dom0 whenever an information on window creation in AppVM is received from qubes_gui.
- Whenever the local window receives XEvent, pass information on it to AppVM (particularly, mouse and keyboard data).
- Whenever AppVM signals damage event, tell local Xorg server to repaint a given window fragment.
- Receive information about window size/position change, apply them to the local window.



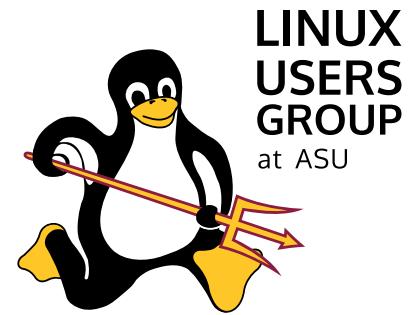
Updating Windows

1. qubes_gui

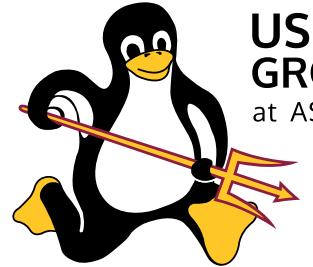
1. Asks the qubes_drv driver for a list of physical memory frames that hold the composition buffer of a window.
2. Passes on this information to the qubes_guid driver in dom0 via a MFNDUMP message.

2. qubes_guid

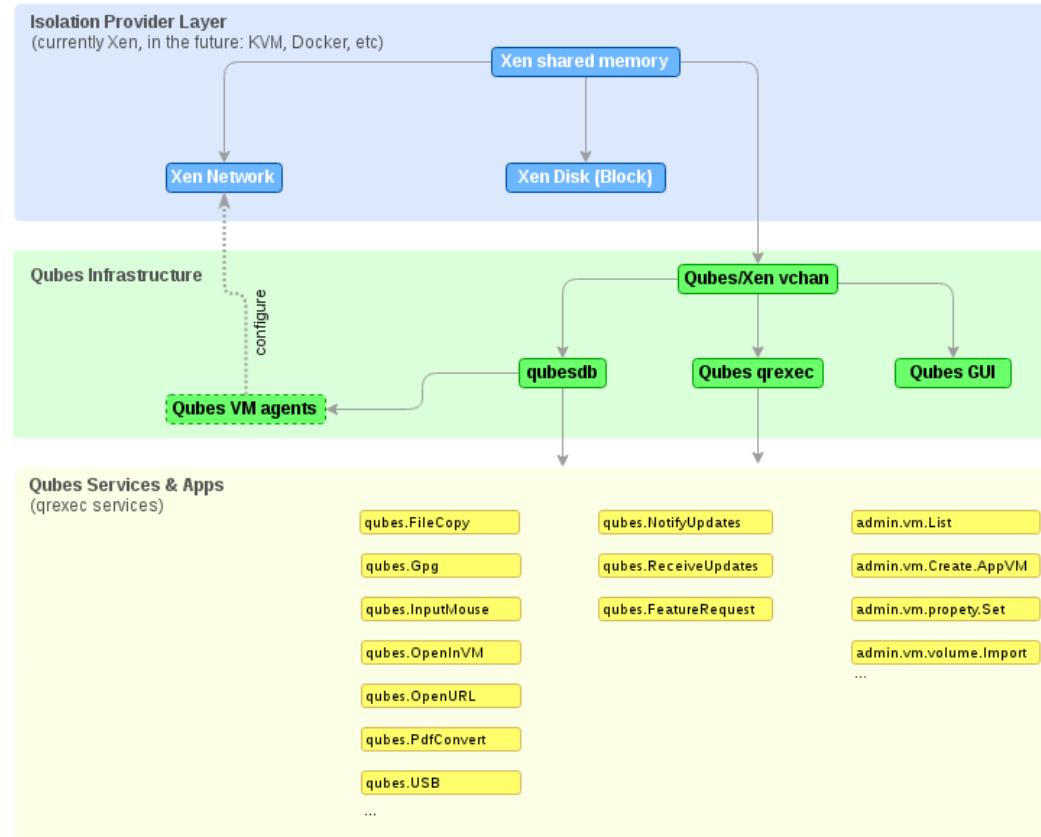
1. Xorg server is started with a LD_PRELOAD-ed library named shmoveoverride.so which hooks all functions calls related to shared memory.
2. qubes_guid creates a shared memory segment and then instructs Xorg to attach it via the MIT-SHM extension.
3. When Xorg tries to attach the segment via glibc shmat, shmoveoverride.so intercepts the calls and instead maps AppVM memory via xc_foreign_pagers
4. We can then use MIT-SHM functions to draw onto a dom0 window.

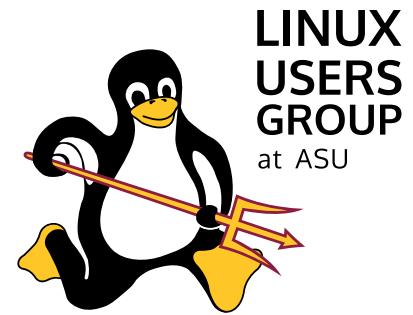


Section 4.4 – General Overview

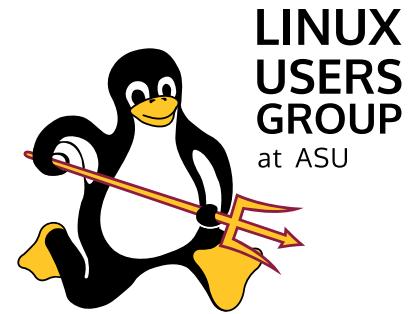


Overview



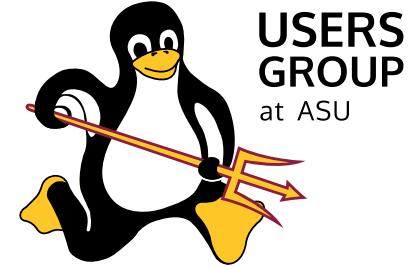


Section 5 – Security and Attack Surfaces

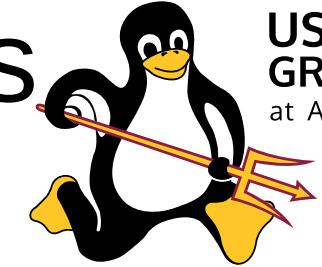


Section 5.1 – Project Security

Security Team

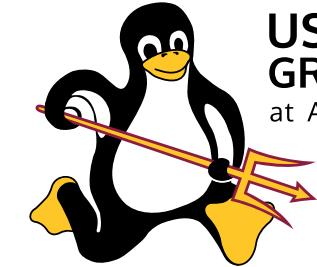


- Handles disclosed security issues.
- Evaluates the importance of security advisories of dependencies.
- Applying security patches.
- Publishing security bulletins and canaries.
- Managing the project's PGP keys.



Security Team Members

- Marek Marczykowski-Górecki – Project Lead
- Simon Gaiser – Project Developer
- Joanna Rutkowska (canaries only) – Project founder



Signing and PGP Keys

- Qubes has signing keys for each major release.
- Every Qubes team member's PGP keys are published on the Qubes key site.
- Every member of the security team must sign a security release.

Canaries

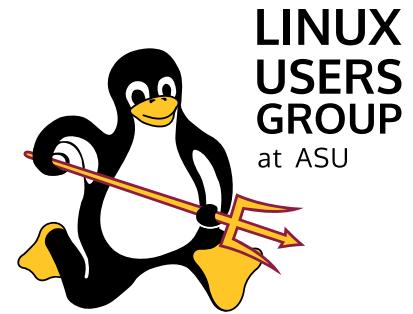


- Canaries are periodical security announcements by the Qubes security team.
- Requires every member sign them.
- Includes specific information such as the date, its number, when the next one plans to be, etc.
- Includes proof of time by news articles and the hash of a Bitcoin block mined in the last 24 hours before the release.

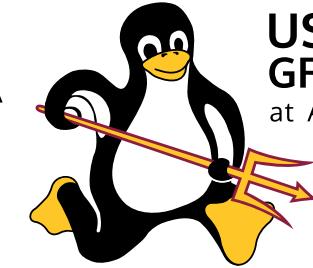
Security Bulletins



- Special announcement issued by the security team.
- Contains summary and analysis of a specific set of vulnerabilities and their patch.

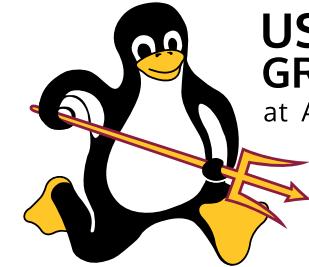


Section 5.2 – Hardware Mitigation



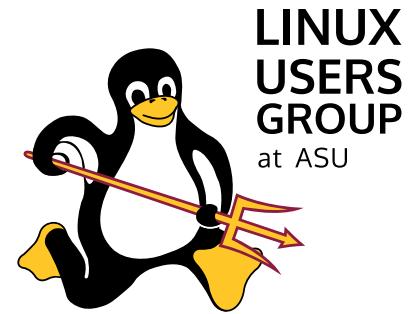
Disk Encryption and MFA

- LUKS 2 full disk encryption.
- U2F Authentication – Only an alternative
- Multi-factor Authentication – Only for restricting access to external services, not the login.

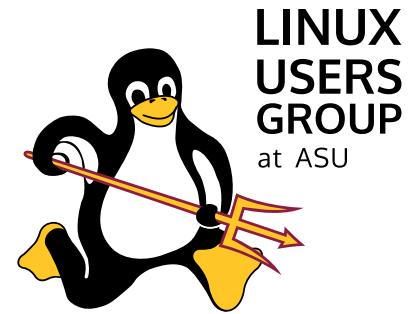


Anti Evil Maid

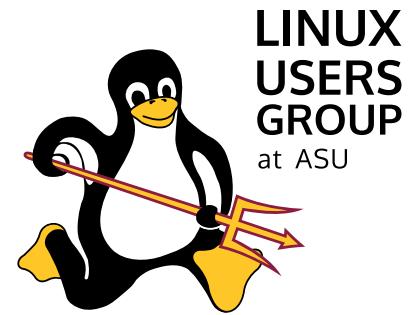
- External TPM trusted boot.
- If the whole chain is valid and properly executed, a message will be loaded validating it.
- Requires USB devices to be able to attach directly to dom0.



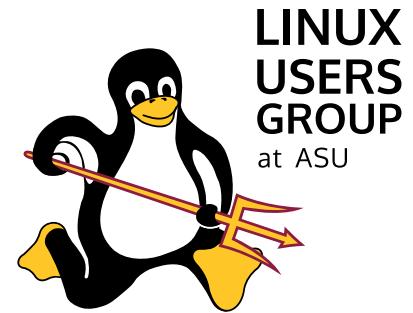
Section 5.3 – Attack Surfaces and Risk



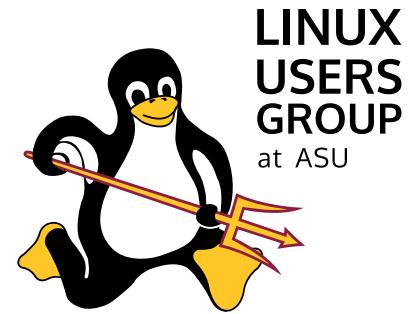
Physical Security of the Device



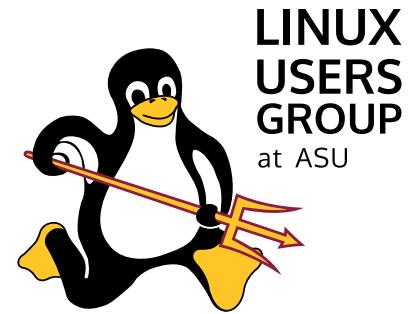
Hypervisor



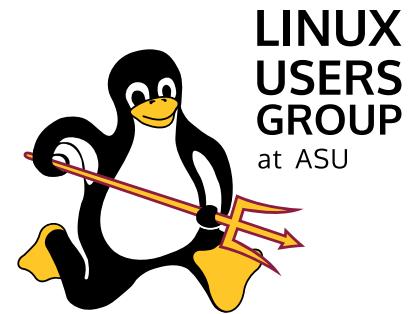
dom0 Drivers



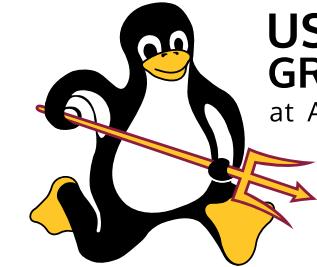
Custom-built Malware



Real World Risk?



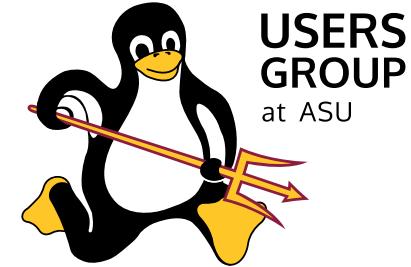
Conclusion



Check out the Project

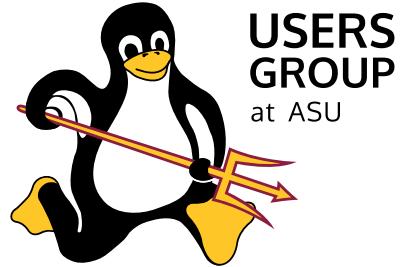
- QubesOS official website: <https://www.qubes-os.org/>
- Download QubesOS:
<https://www.qubes-os.org/downloads/>
- QubesOS Forums: <https://forum.qubes-os.org/>
- QubesOS GitHub: <https://github.com/QubesOS>

Support the Project

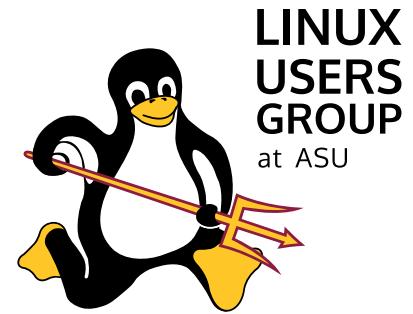


- QubesOS Donation:
<https://www.qubes-os.org/donate/>
- Contribute:
<https://www.qubes-os.org/doc/contributing/>

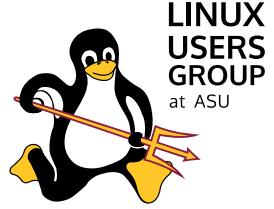
Resources



- Qubes Website: <https://www.qubes-os.org/>
- Qubes Documentation: <https://www.qubes-os.org/doc/>
- Qubes GitHub: <https://github.com/QubesOS>
- Xen Project: <https://xenproject.org/>
- Xen Project Wiki: https://wiki.xenproject.org/wiki/Main_Page
- Invisible Things Lab: <https://invisiblethingslab.com>
- Tarsnap website and documentation: <https://www.tarsnap.com/>



Questions?

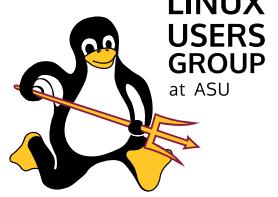


Qubes OS

A Technical Introduction

Daylam Tayari

Technical?

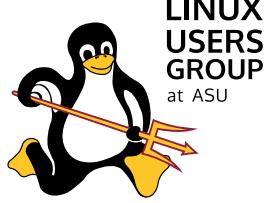


Don't worry – This won't be exclusively technical

Overview → Specificities → Technical Details

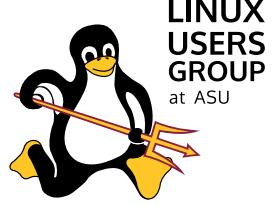
Questions at any point?
Ask!

Talk Layout



1. Overview
2. Usage and Security Implementations
3. Drivers, Interfaces and Policies
4. Technical Analysis of the Architecture
5. Security and Attack Surfaces

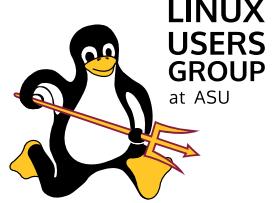
© 2022 Daylam Tayari - CC BY-NC-SA 4.0



Section 1 – Overview

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

What is Qubes?



- An operating system based on a hypervisor.
- Everything compartmentalised.
- Separation at every level – Hardware does not interact directly with the end usage (normally).
- All your regular actions are in VMs.
- Entirely open source.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

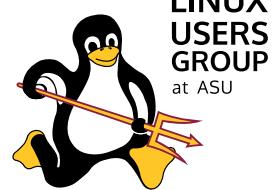
Based on the Xen hypervisor.

Compartementalised at every level – The core, the hardware facing VMs, template and app VMs, disposable VMs.

Why normally? PCI devices can still be passed through.

All the work is done in the AppVMs.

History



- Introduced in January 2010.
- Founded by Joanna Rutkowska.
- Compartementalised approach for security.
- Currently led by Marek Marczykowski-Górecki.
- Latest version is 4.1.0 released 4th of February 2021.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

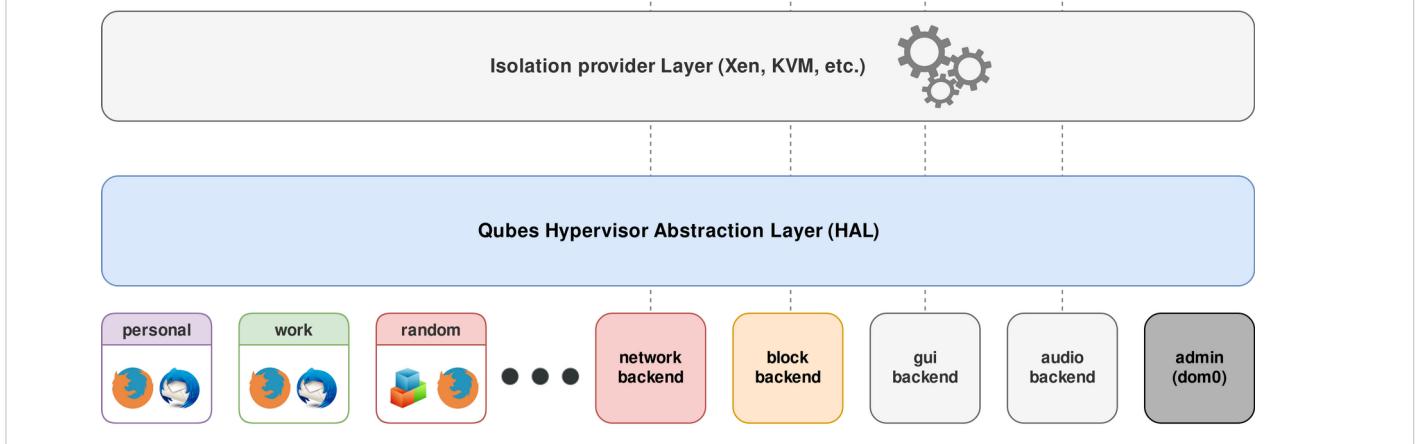
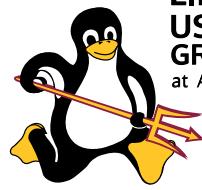
Rafal Wojtczuk was part of the initial release.

Through Invisible Things Lab.

Work on it had started a couple months prior.

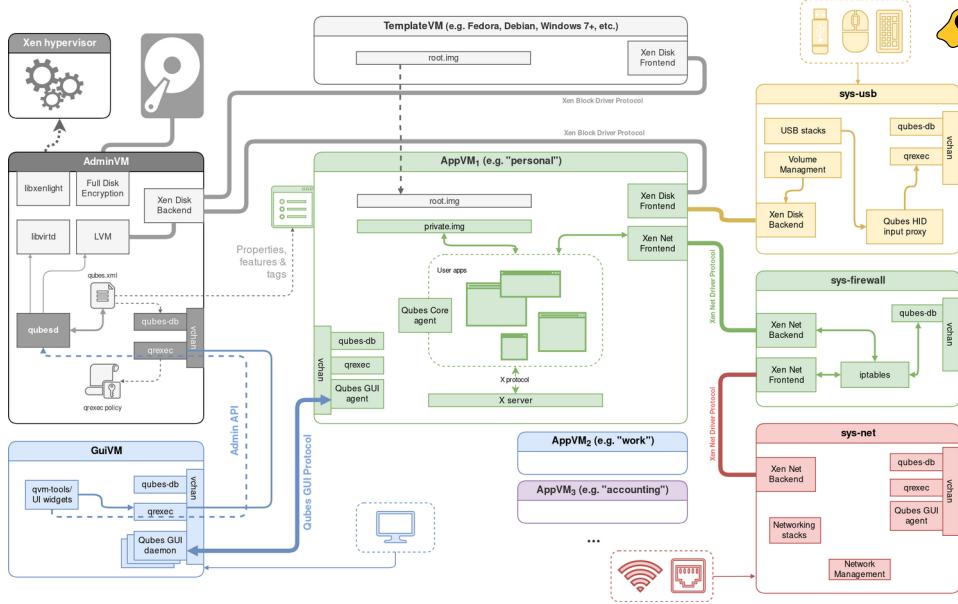
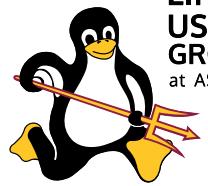
Architecture

LINUX
USERS
GROUP
at ASU



Core Stack

LINUX
USERS
GROUP
at ASU



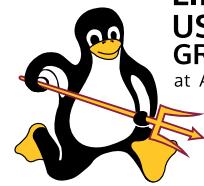
Working in AppVMs



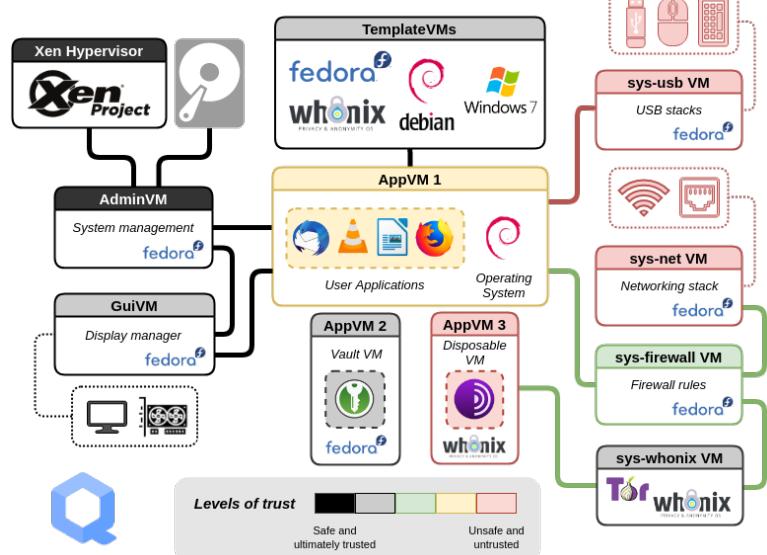
- AppVMs are where you do all your work.
- They can run everything – even Windows
- Connected to the hardware through other VMs.
- Interact with the core and other VMs with drivers.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Mention how I will often refer to VMs as a Qube or as Qubes.

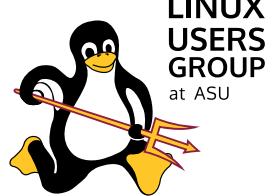


Use Case Example



© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Why?



- Security by Isolation Approach
 - Virtualisation
 - Compartementalisation.
 - Need to Have approach.
 - Disposable VMs.
- Privacy
 - Whonix and Tor
 - Preventing Fingerprinting.
- Compartamentalisation of your Workspace.
- Changing Systems Often

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

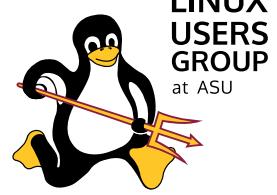
Need to have – Only give VMs what they need. Like need to know.

Fingerprinting – Just over 40,000 people use Qubes to any capacity. (fingerprinting based off update server connections: <https://www.qubes-os.org/statistics/>)

Like compartmentalising everything (personal anecdote)

Changing systems often – Cross system migration with Qubes backups. Also practicality of temporary concealment (border crossing example – privacy rights, i.e. Australia but also the US).

Why Not?



- Technical experience occasionally required.
- Resources
 - Overhead
 - Sharing
 - Device Specifications
- Thinking in a Qubes way
- Technical Restrictions

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Technical experience includes being able to troubleshoot, know what something means to pinpoint the error and go from there.

Overhead around 4 Gb of RAM and very low CPU wise for Dom0. Resources are shared between VMs but can also be limited. Also means each VM will have some level of overhead.

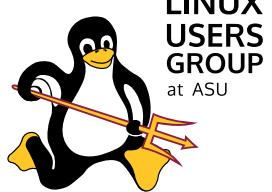
Requirements:

- 64 bit Intel or AMD processor with virtualisation.
 - Intel: VT-x with EPT and VTT-d
 - AMD: AMD-V with RVI and AMD-Vi/AMD IOMMU
- The latter is I/O MMU (Memory Management Unit) virtualisation which allows VMs to access PCI devices (PCI passthrough).**
- 6Gb of RAM (16gb recommended)
- 32Gb of storage (128Gb recommended)

Thinking in the Qubes way – Thinking in the way of compartmentalising everything not just using a central VM or two. Technical restrictions – Screensharing, full size screenshots, USB devices can only be passthrough a single VM at a time (i.e. audio output or mic).

Certain customisation which include Dom0 can risk both security and be a technical challenge.

Why do I?



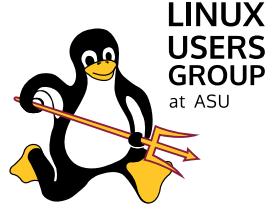
- Compartementalisation
- Security
 - Isolation
 - Need to Have
- Mirroring across devices
- Whonix

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Always compartmentalised growing up.

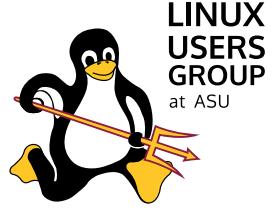
Using various operating systems for different uses.

Compartementalisation also at the network level: Ethernet, WiFi, VPN, Tor, etc.



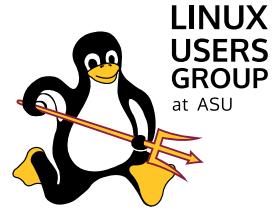
Section 2 – Usage and Security Implementations

© 2022 Daylam Tayari - CC BY-NC-SA 4.0



Section 2.1 – Virtual Machines in Qubes

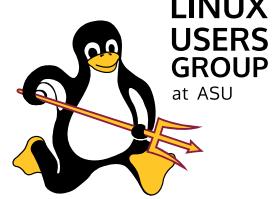
© 2022 Daylam Tayari - CC BY-NC-SA 4.0



Section 2.1.1 – Virtual Machine Types

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Types of VMs

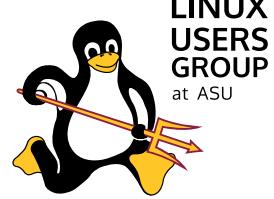


- DispVMs – Based off template and no persistance.
- AppVMs – Based off a template and limited persistance.
- TemplateVMs – Root file system for AppVMs.
- StandaloneVMs – ‘Traditional’ VM.
- AdminVM – The core, dom0.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Everything is essentially a VM but which
are for what.

DispVMs



- Disposable VMs are temporary disposable systems.
- Based off a template.
- Inherits the entire template including `/home/`.
- No persistence whatsoever, fully wiped on exit.
- Has a temporary volume though can be set to only use RAM.
- Files can be easily opened to be viewed.
- Randomly numbered between [1-10000).
- **Demo!**

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Each DispVM uses volatile.img file that is used to create R/W illusion for its root fs via COW (same mechanism as for any other AppVM based on a template). This file is located in /var/lib/qubes/appvms/<name-of-dispvm-template>/volatile.img.

(BTW, an interesting peculiarity, which you, like myself, might be wondering about: when another DispVM is started this volatile.img, used by the previous DispVM is unlinked, and so the new DispVM uses a new copy of the file, inode-wise).

Theoretically, one could modify the conf file used for dispvm creation, specifically this line (e.g. in /var/lib/qubes/appvms/fedora-18-x64-dvm/dvm.conf):
'script:file:/var/lib/qubes/appvms/fedora-18-x64-dvm/volatile.img,xvdc,w',

... and change the path to point to /dev/shm/qubes, where the DispVM savefile is kept, and which is stored in RAM entirely (the whole /dev/shm is a ram-based fs).

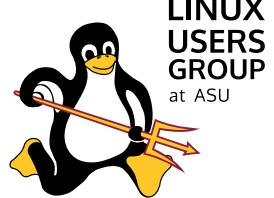
But, as Marek pointed out, a DispVM which might go to the dark side of the force, might then generate writes to its whole root fs and make this volatile.img file to swallow up to 10GB in size, making the whole system unusable if we kept it in /dev/shm.

Unless we could setup some quota on /dev/shm, can we?

Anyway, regarding your forensics concerns -- a rather easy and convenient way to remove any records of a particular AppVM's activities is to simply shred its private and volatile images:

shred private.img volatile.img

AppVMs

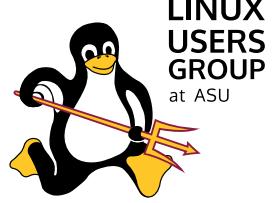


- Based off a template.
- Entire `/rw/` directory persists which includes:
 - `/home/`
 - `/usr/local/`
- Inherits part of its `/home` and `/usr/local`.
- Is provided an image of the template VM which was **generated at its last shutdown**.
- Where you will do all your work.
- **Demo!**

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Inheritance from template includes the `/etc/skel` directory which points to `/home/` and the `/usr/local.orig/` directories in the template.

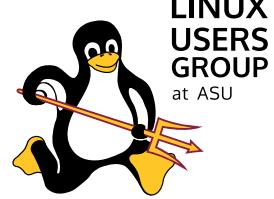
Template VMs



- Templates have no inheritance, recursive templates is not a possibility.
- Everything persists so **be careful**.
- No network by default.
- Its image is updated at shutdown.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Standalone VMs



- Traditional VMs.
- Either independent clone of template or separately installed system.
- No integration though possible with drivers.

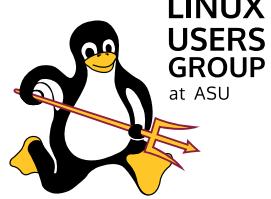
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

You can install a VM using an ISO.

By adding the Qubes drivers to the system, it can have integration with the rest of the system.

Qubes Windows Tools.

AdminVM



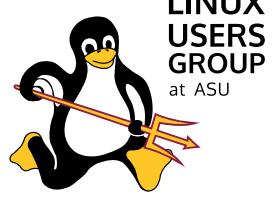
- AdminVM is dom0.
- Is the backbone of the operating system.
- Where the entire backend of the OS is ran.
- No network connectivity.
- Updates can only be made to specific repositories.
- Runs Fedora 32.
- Isolated from the rest of the system – No clipboard integration.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

GUI VM is planned to be moved out into its separate VM in a future release.

Explain logic behind Fedora 32.

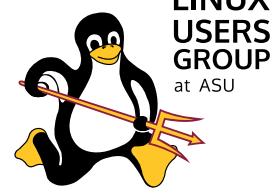
GUI VM is in its own Qubes.



Section 2.2 – Template VMs

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Template Repositories

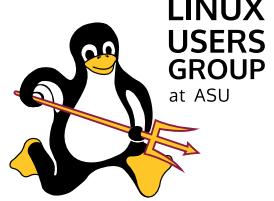


- Official Templates – Built and released by the Qubes team.
- Community Templates – Designed, built and maintained by the community.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Community Templates are only those that are accepted and added to the community template repository, not just any.

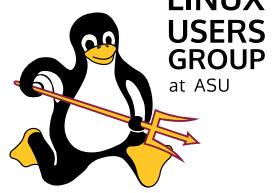
Official Templates



- Fedora 34
 - Minimal Template
 - Xfce Template using Xfce instead of GNOME.
- Debian 11
 - Minimal Template
- Whonix 16 Gateway and Workstation

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

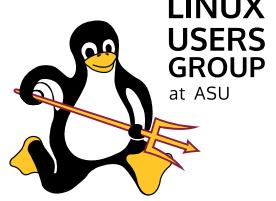
Community Templates



- Whonix
- Ubuntu
- Arch Linux
- CentOS
 - Minimal Template
- Gentoo
 - Minimal Template
- NetBSD

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Custom Templates



- Templates which aren't standalone but can be implemented.
- Kali Linux
- Parrot

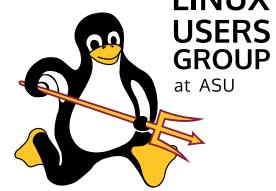
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Updating



- Qubes has an updating utility which will handle proper updating for Templates.
 - Creates a mirror.
 - Updates.
 - If successful → Updates the actual Template
- Manually in the individual TemplateVM.

Installing Programs

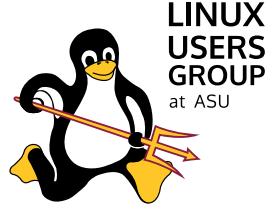


- Using the package managers.
- Sending over binaries and repositories and installing them locally.
- Adding them manually in directories such as `/opt/` and creating respective desktop files.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

No internet access in Templates by default and they shouldn't be granted it.

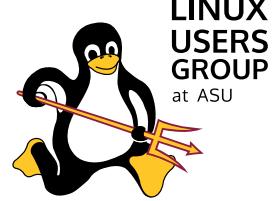
Whenever the contents of the elevated clipboard are retrieved and sent to a VM, they are wiped from the elevated clipboard.



Section 2.3 – Migration and Backups

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

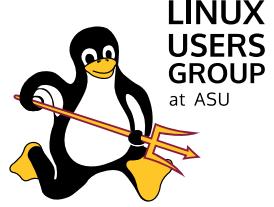
Cloning



- Any VM, regardless of type can be easily cloned.
- Cloning a VM is accessible directly in the basic settings.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

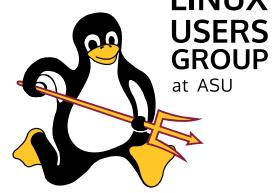
Qubes Backups



- Qubes official backup utility.
- Packages a set of VMs into a backup.
- Saves the version of the last shutdown.
- Saves specific file systems:
 - AppVMs – private.img
 - TemplateVMs – root.img
 - StandaloneVMs – root.img and private.img
- Encrypts backups with scrypt.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

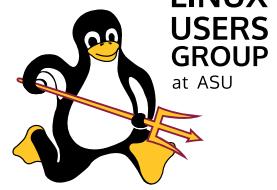
Backup Encryption



- Uses the scrypt encryption utility.
- Initially developed for the Tarsnap online backup system.
- Uses a complex cryptographic system including both signing and encryption, both which are required for decryption.
- Qubes uses their release signing key for the signing.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Tarsnap Cryptography

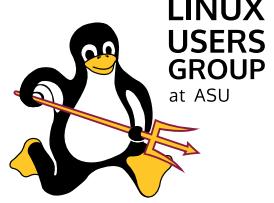


- A 2048-bit RSA key used for signing archives. This is used in combination with SHA256 and a Merkle hash tree to verify the authenticity of stored archives.
 - Cryptographers: RSASSA-PSS is used, with SHA256 as the hash function.
- A 2048-bit RSA key used for encrypting session keys. All the data which the Tarsnap client sends to the server to store is encrypted with per-archive random AES-256 keys; those keys are encrypted with this RSA key and attached to the stored data.
 - Cryptographers: RSAES-OAEP is used to encrypt the session keys, using MGF1 and SHA256, and the padding verification performed when decrypting is carefully written to be free of timing side channels. The AES-256 keys are used in CTR mode, with sequentially incrementing nonces — generating a new AES-256 key for each session ensures that a key-nonce pair will never be used twice.
- A 256-bit HMAC-SHA256 key used to protect each individual block of data from tampering. From a cryptographic perspective, this is unnecessary, since a Merkle hash tree protects each archive; but data is compressed using zlib before being stored, so this provides protection against a theoretical attacker who can tamper with stored data and has found a security flaw in zlib decoding.
 - The zlib decoding process is quite complex — for that matter, the same is true of all decompression algorithms — and such complex code has a relatively large risk of having security vulnerabilities, hence the decision to add an extra layer of security here.
- Two 256-bit HMAC-SHA256 keys used to generate names for blocks of data stored. Tarsnap uses the same reference-by-hash trick as the author's Portsnap and FreeBSD Update utilities; using HMACs instead of raw SHA256 hashes prevents any information from leaking via the hashes.
 - Why two keys? One is used to hash data, and the other is used to hash archive names. Yes, that's right, even the names of archives are stored securely!

<https://www.tarsnap.com/crypto.html>

©2022 Daylam Tayari - CC BY-NC-SA 4.0

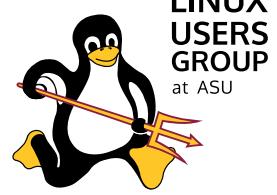
Backup Restore



- Official Qubes Restore utility.
- Restores backups without a hitch.
- Has the possibility to solely verify integrity.
- Can either restore alongside or overwrite.
- Restoration has no difference machine to machine as long as they are on the same Qubes release.

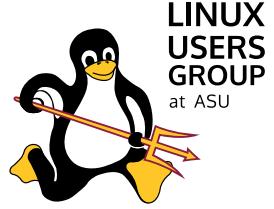
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Migration to an ISO



- Official support with Qubes Builder
- Allows you to turn a Qubes image into an ISO
- Theoretically should be flashable onto another computer just like any ISO.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

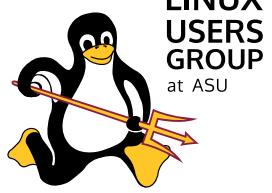


Section 2.4 – Interactions between VMs

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

This is a short overview of the various interfaces and drivers.
We will go into further detail in section 3.

Clipboard Interactions

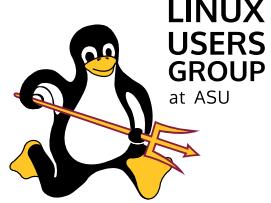


- A single elevated clipboard.
- Any VM can send their clipboard to the elevated clipboard with `Ctrl+Shift+c`.
- Any VM can retrieve the contents of the elevated clipboard with `Ctrl+Shift+v`.
- Elevated clipboard is wiped on retrieval.
- AdminVM is not integrated with the elevated clipboard.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Explain why dom0 pasting is prohibited.
Security.

File Interactions

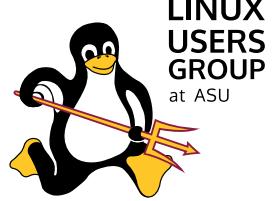


- All VMs can send files to another VM.
- All VMs except the AdminVM can receive files.
- Files will arrive in the `~/QubesIncoming/VM_NAME/` folder.
- AdminVM can only receive files from other VMs by passing the IO from the VM to the AdminVM.
 - i.e. `qvm-run --pass-io QUBE_NAME 'cat NAME_OF_FILE' > NAME_IN_DOM0`

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

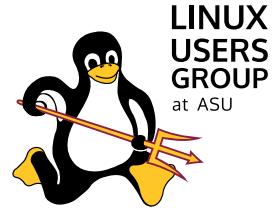
Explain why dom0 passing is so complicated.

Split Utilities



- Split GPG
- Split SSH
- Split Electrum Wallet
- Split dm-crypt
- U2F Proxy

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

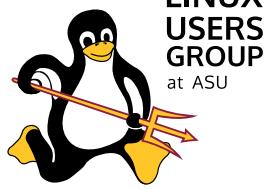


Section 2.5 – Networking

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

This is a short overview of the various interfaces and drivers.
We will go into further detail in section 3.

Networking Overview



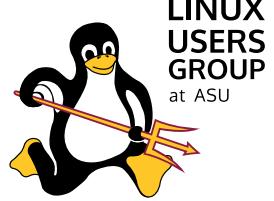
- Any VM can be set to be a networking VM.
- There are two primary VMs responsible:
 - Hardware VM (i.e. WiFi, Ethernet, etc.)
 - Firewall VM
- Other VMs can be configured to handle and adapt a network connection.
- Network VMs can be chained.
- Changing networks is as simple as changing one setting and is applied instantly.
- You select the final VM as the one you want to use.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Highlight the chained part and elaborate a bit on it.

Explain the final selection and bring up the VPN example.

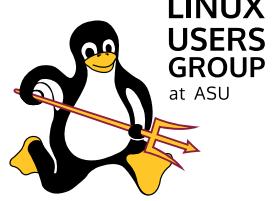
Hardware VMs



- Hardware networking VMs are those that interact with the actual PCI devices.
- Good practice is one per PCI device.
- Lets everything through.
- Uses NetworkManager by default.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Firewall VM



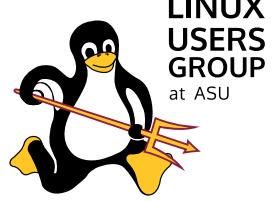
- Firewall rules can be set for each Qube.
- Applies all network-level policies.
- Responsible for blocking inter-VM networking.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Firewall policy can be from allow all/allow none to allow only specific whitelist IPs and only to specific ports and protocols.

Inter-VM networking can be blocked.

Custom VMs

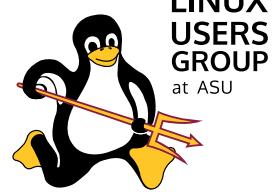


- No limit in the amount of network VMs that can be created and made.
- You can create a VPN VM and everything that is passed through goes through the VPN.
- You can deploy a full on DNS sinkhole (i.e. Pi-hole) or a firewall if you wanted and chain it.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Bring up the chaining aspect and chaining through multiple VMs.

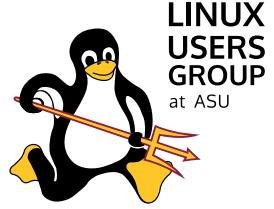
Whonix Networking



- sys-whonix VM running Whonix Gateway.
- All connections passing through it are passed through Tor.
- Tor connection is fully configurable.
- Just like any other VM.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Mention creating multiple sys-whonix for multiple connections.

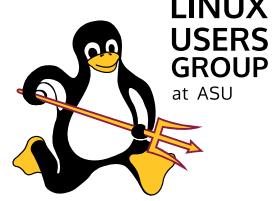


Section 2.6 – USB Handling

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

This is a short overview of the various interfaces and drivers.
We will go into further detail in section 3.

USB Qubes



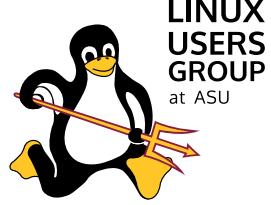
- Separate VM which interacts directly with the USB PCI device.
- Without a USB Qube, the USB PCI device interacts directly with dom0.
- To use or have access to a USB device in a VM, it HAS to be manually passed through the Qubes Devices taskbar menu.
- **Demo!**

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Passing as a USB keyboard security risk.

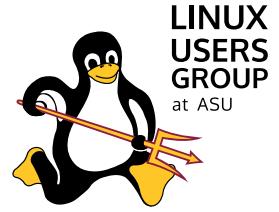
Camera example.

Complications



- Necessary passthrough for login
 - Passthrough of USB keyboard device to dom0 on boot through a policy.
- Connection not as stable as if directly passed through.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

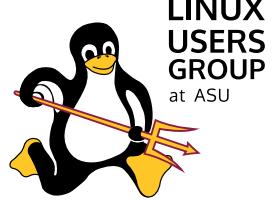


Section 2.7 – VM Configuration

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

This is a short overview of the various interfaces and drivers.
We will go into further detail in section 3.

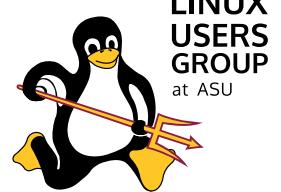
Basic Settings



- Name and label colour
- Template
- Networking VM
- Private storage size
- Start on boot
- Include in backups
- Debug mode

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Advanced Settings



- Initial Memory
- Max Memory (only if memory balancing is enabled)
- VCPUs Number
- Kernel
- Virtualisation
- Disposable VM Template (if not itself)
- Provides Network Toggle
- Boot from ISO

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Memory balancing is also present

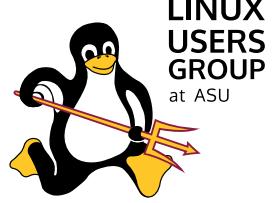
Kernel Options:

- 5.4.136
- 5.4.143
- 5.4.156
- 5.10 (Qubes 4.1)

Virtualisation options:

- PVH
- HVM
- PV

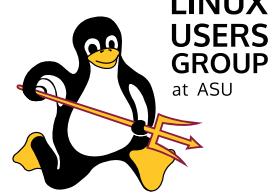
Firewall Rules



- Configure firewall rules.
- Configurable by:
 - Address
 - Port
 - Protocol
- Temporary Access

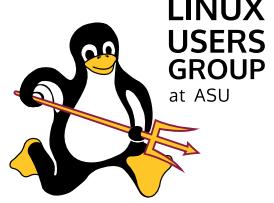
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Devices



- PCI Passthrough
- Must use HVM or PV virtualisation and have disabled memory balancing for ideal results.
- Limit in types of PCI device that can be passed depending on the VM type.

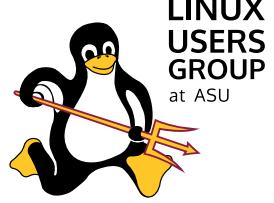
Applications



- Select which applications show up in the application menu for that VM.
- Grabbed from all `*.desktop` files in the `/usr/share/applications/` directory.
- Refreshed from Template VM.

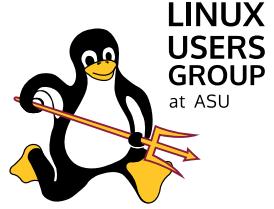
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Services



- Add or remove specific services.
- Enable or disable specific added services.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

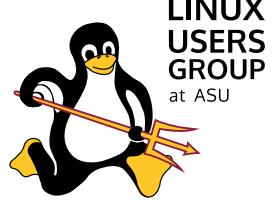


Section 2.8 – Customisation

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

This is a short overview of the various interfaces and drivers.
We will go into further detail in section 3.

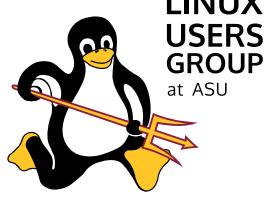
DEs and WMs



- Desktop Environments officially supported:
 - Xfce (default)
 - KDE
- Window Managers (community support):
 - i3
 - AwesomeWM
 - bspwm (unofficial)

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

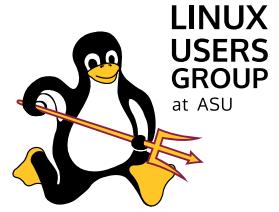
Security when Ricing



- Passing programs to dom0
- Affecting dom0
- Restriction in availability of programs.

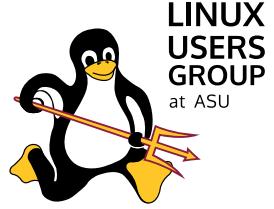
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

You can't go around and add any old
program to your dom0.
Also maintaining them and security risk.



Section 3 – Drivers, Interfaces and Policies

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

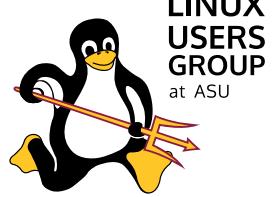


Section 3.1 – Intercommunication Drivers

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

This is a short overview of the various interfaces and drivers.
We will go into further detail in section 3.

Clipboard

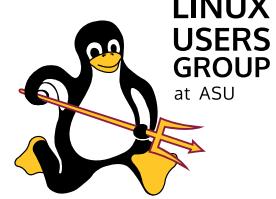


- A single elevated clipboard.
- Any VM except dom0 can copy and paste from it.
- It can only store a maximum of one entry.
- When it is pasted in a VM, it is wiped.

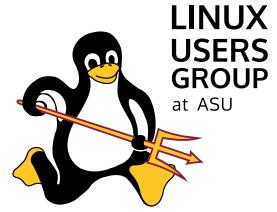
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Can also be manually cleared.

File Transfer



- Files can be transferred between all VMs except dom0.
- Files will be sent to the `~/QubesIncoming/VM_NAME/` directory.

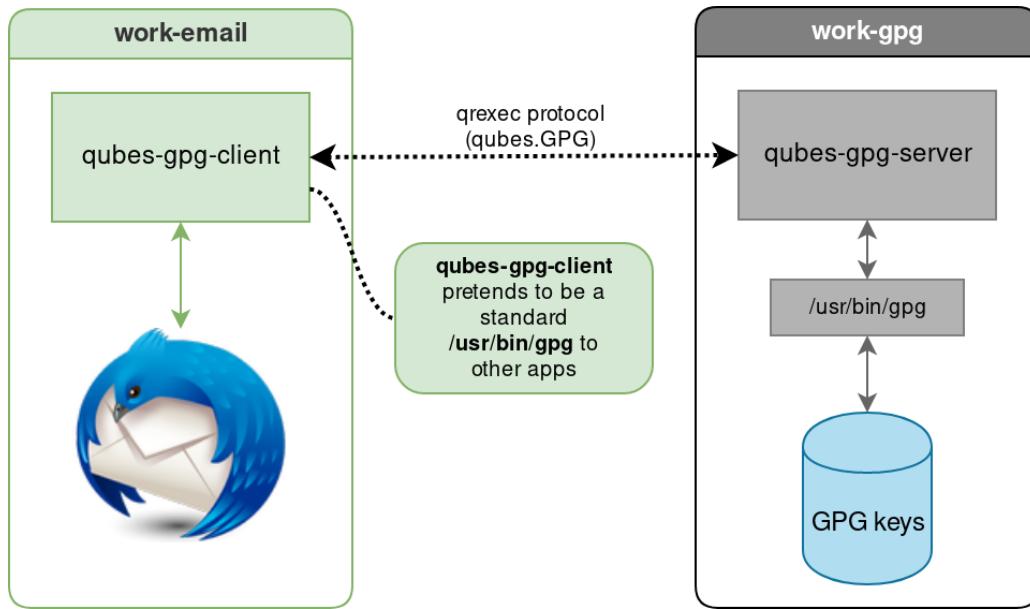
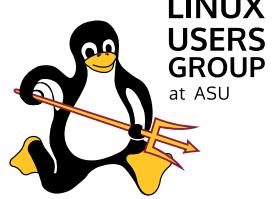


Section 3.2 – Split Interfaces

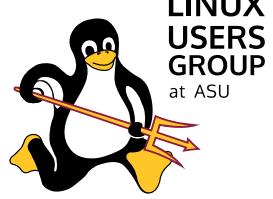
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

This is a short overview of the various interfaces and drivers.
We will go into further detail in section 3.

Split GPG - 1



Split GPG - 2



- All keys are stored in a separate VM which is ideally fully isolated.
- The qubes-gpg-client is what interacts with the client, passing all the commands over to the GPG program on the GPG Qube and sending its output back.
- Two sets of prompts appear at each GPG usage.
- Access can be determined to a set amount of seconds.
- Passwords are entered and validated directly in the GPG Qube.
- You set your GPG client to be the Qubes client in programs.
- **Demo!**

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Isolation = no internet connection or devices or anything ever passed to it (such as USB devices)

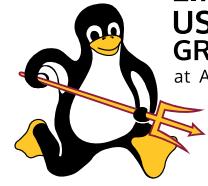
The policy is determined in the `'/rw/config/gpg-split-domain` file.

The password for PGP key never touches an AppVM.

Example for like Thunderbird and Git.

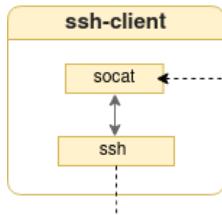
Split SSH - 1

LINUX
USERS
GROUP
at ASU



Plain Setup

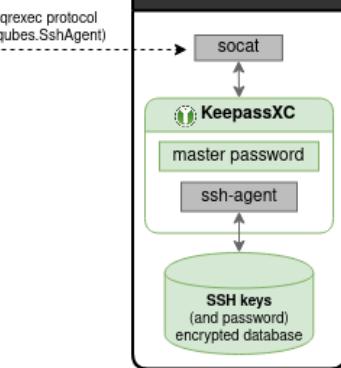
uses ssh-askpass



Remote Server

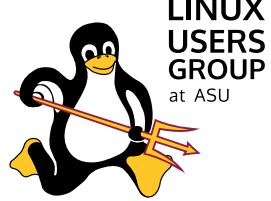
KeePassXC Setup

Uses the password manager to protect your keys



© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Split SSH - 2

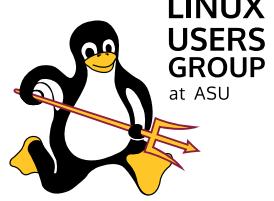


- SSH keys safe in an isolated VM.
- SSH keys are retrieved and passed through when requested.
- SSH private keys are kept in the isolated VM.
- Unfortunately, use can be held on for the entire of the existence of the VM.
- Community made.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

SSH agent connection, once done, can be kept active and held on for the entirety of the time the VM stays up which is a big liability and my personal reason why I don't use it.

Split dm-crypt

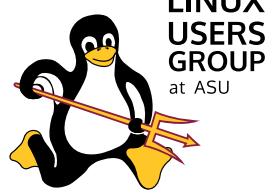


- Splits up the process of decryption of a LUKS1 partition.
- Encrypted VM is attached to a specifically configured offline disposable VM.
- Prompts for a password and decrypts the partition in the disposable VM.
- Mounts the partition into the destination VM.
- The destination VM does not have access to the password or encryption key.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

The encrypted partition is attached from the source VM to an offline device DisposableVM configured not to parse its content in any way: The kernel partition scanners, udev probes, and UDisks handling are disabled.

Split Electrum Wallet

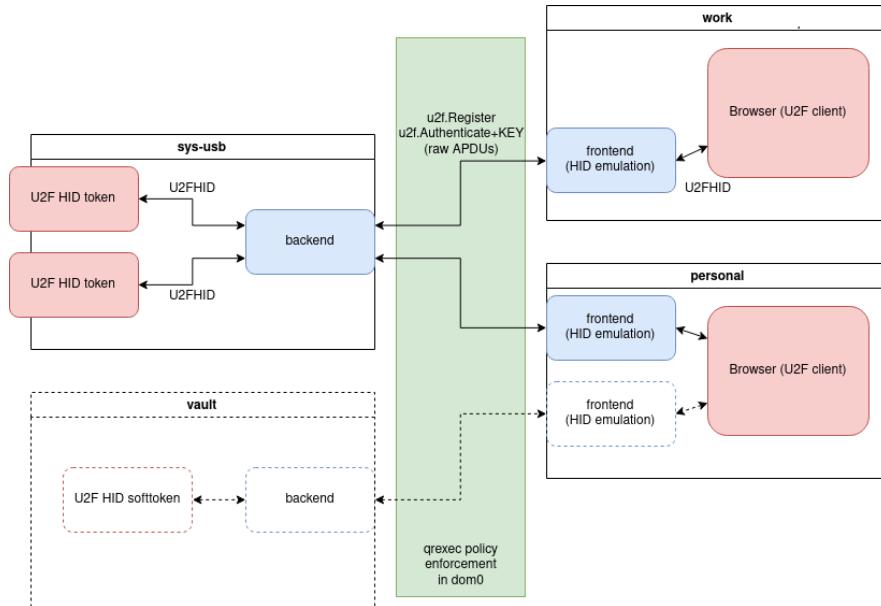
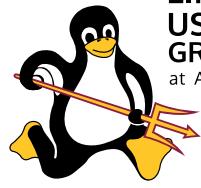


- Wallet is stored in one isolated offline VM.
- Wallet is accessed via an online watching-only Electrum wallet which connects to the isolated wallet.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

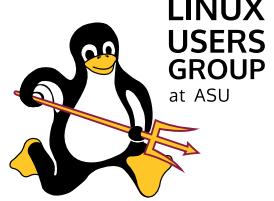
U2F Proxy - 1

LINUX
USERS
GROUP
at ASU



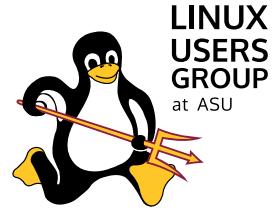
Dashed lines signify upcoming
enhancements.
Soft tokens to emulate a smart card.

U2F Proxy - 2



- The proxy only passes the necessary to the actual U2F key, performs the transaction and returns it to the browser without passing the entire U2F key to the device.
- Can be set to only allow specific tokens per VM.
- No direct interaction between the U2F device and the VM that is requesting the token.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

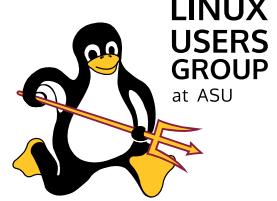


Section 3.4 – Policies

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

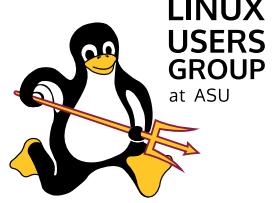
This is a short overview of the various interfaces and drivers.
We will go into further detail in section 3.

Policies in Qubes



- Each interaction with a driver is controlled by a policy.
- Policies can be set both inside of a VM for its drivers or for the system.
 - `/rw/config/`
 - `/etc/qubes-rpc/policy/` in dom0

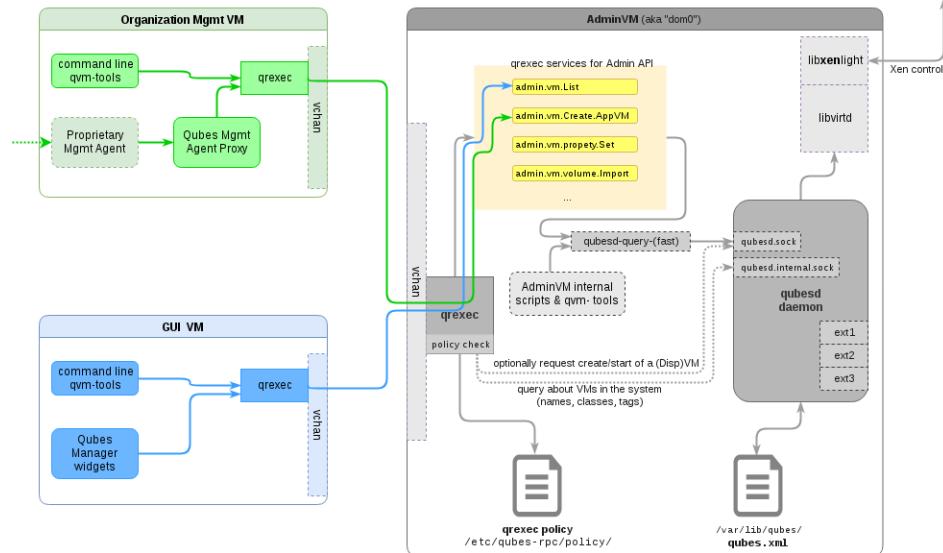
Control of Policies



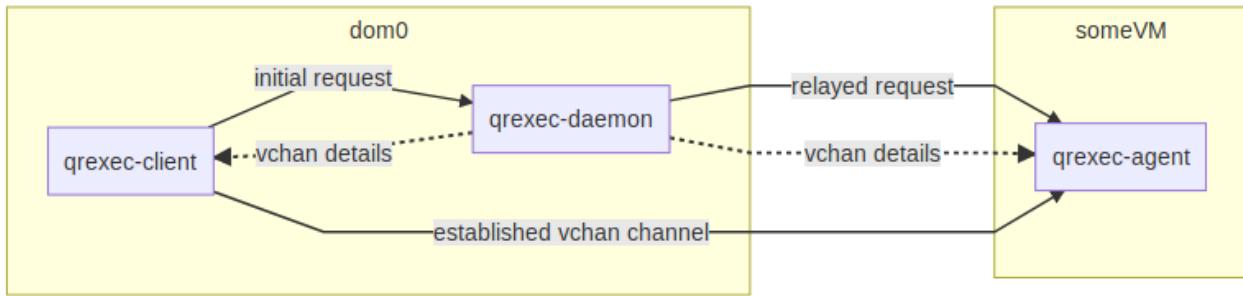
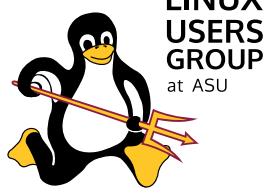
- Which VMs does the policy apply to.
- How to handle any VM not explicitly referred to.
- Handle policies depending on the tag of a VM.
- Allow or deny specific actions.

Policy Overview

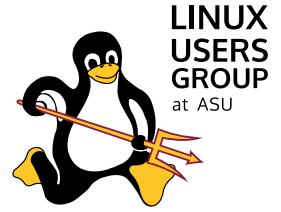
LINUX
USERS
GROUP
at ASU



Policy Client Overview

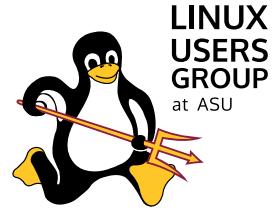


© 2022 Daylam Tayari - CC BY-NC-SA 4.0



Section 4 – Technical Analysis of the Architecture

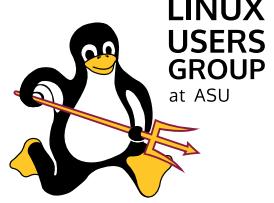
© 2022 Daylam Tayari - CC BY-NC-SA 4.0



Section 4.1 – Volume System

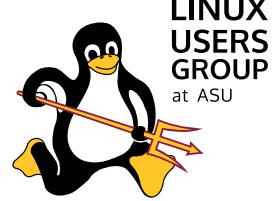
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Volume Types



- root.img – Root device (`/`).
- private.img – Persistant volume of a VM.
- volatile.img – Swap and temporary `/`. Discarded on reboot.
- modules.img – Kernel modules and firmware.

Block Devices

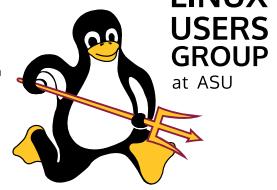


- xvda – Base root (**NOT** root.img)
- xvdb – private.img
- xvdc – volatile.img
- xvdd – modules.img

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

xvda actually represents a snapshot of the root.img volume, not a direct link to it which is why the differentiation.

Root Implementation - 1



- Two versions of `root.img`:
 - `root.img` – True root filesystem.
 - `root-cow.img` – Snapshot of the root filesystem.
- Everytime an AppVM is started, a `root-cow.img` is made (unless no changes).
- AppVMs are based off the `root-cow.img`.

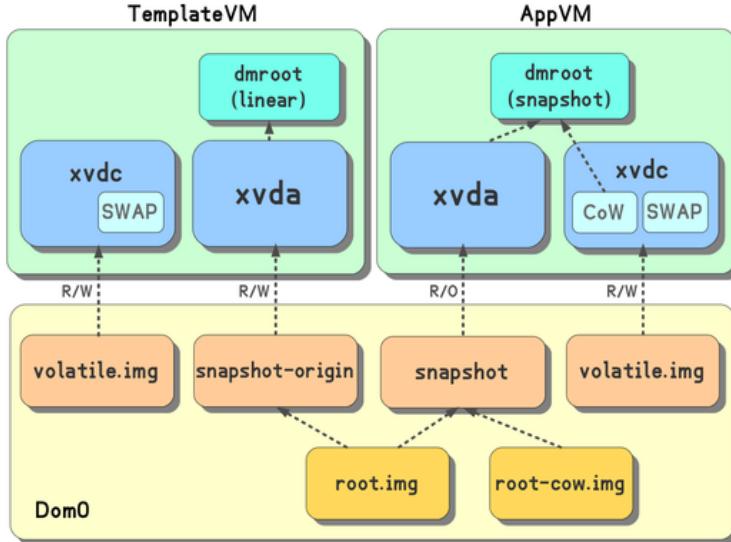
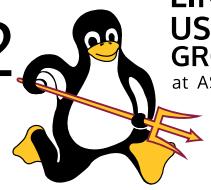
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

CoW stands for Copy on Write

Snapshot made at the start.

If root.img hasn't changed, all the AppVMs will be based on the same template.

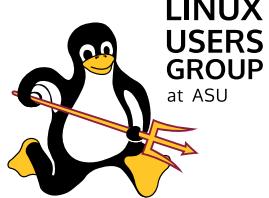
Root Implementation - 2



When a TemplateVM starts, it's provided with the snapshot-orig which it then writes to the root.img during shutdown. Uses the `qvm-template-commit` tool. A Xen script moves the `root-cow.img` volume to ` `.old` and creates a new one with a new inode number. (AppVMs reference to it by inode so no issues and the old root-cow is deleted when the AppVM shuts down).

Reverting to a past root.img IS POSSIBLE.

private.img



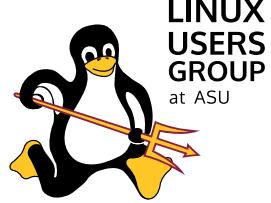
- Persistent storage for an AppVM.
- What is mounted in `/rw/`.
- Mounted in read-write.
- Contains:
 - `/home/` – Bind mounted on `/rw/home/`
 - `/usr/local/` – Symlinked to `/rw/usrlocal/`
 - Qubes configuration files in `/rw/config/` .

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Qubes configuration files such as the PGP.

Show/demonstrate.

volatile.img

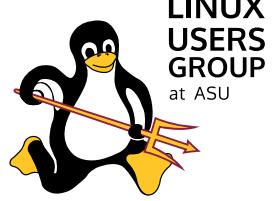


- Contains temporary changes.
- Discarded after AppVM shutdown.
- Contains two primary elements:
 - Changes to the root partition.
 - Swap partition.

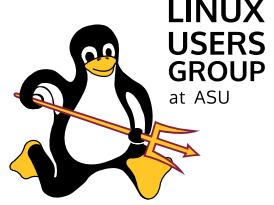
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Changes to the root partition such as installing programs.

TemporaryVMs



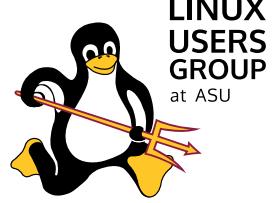
- TemporaryVMs are a unified version.
- Each have their own `root.img` and `private.img` which act in the same way except both are persistent.
- Still use a `volatile.img` for swap.



Section 4.2 – Virtualisation

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

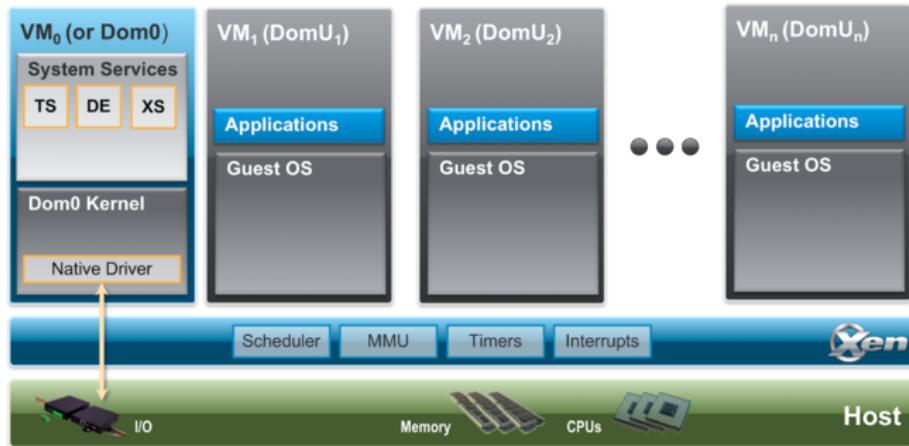
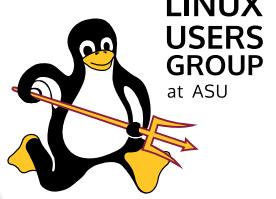
Hypervisor



- Qubes uses the Xen Project Hypervisor.
- Xen is an open source type 1 hypervisor.
- Qubes architecture founded on the Xen Project Architecture

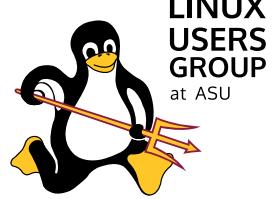
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Xen Architecture



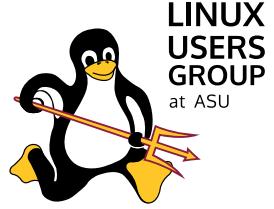
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

VM Virtualisation



- Virtual Machines can be virtualised in three virtualisation types:
 - PVH (default)
 - HVM
 - PV

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

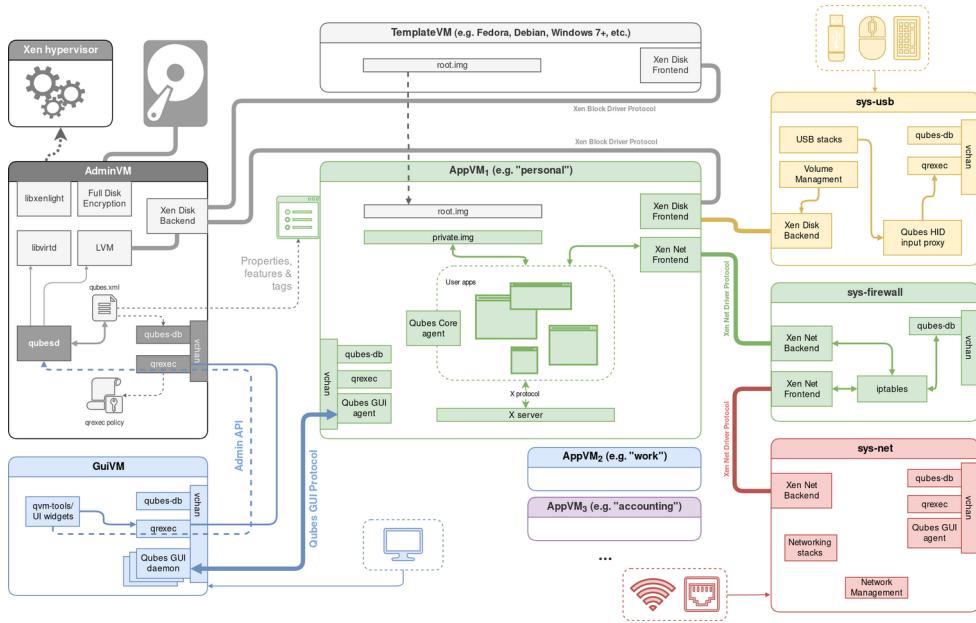
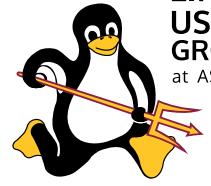


Section 4.3 – GUI

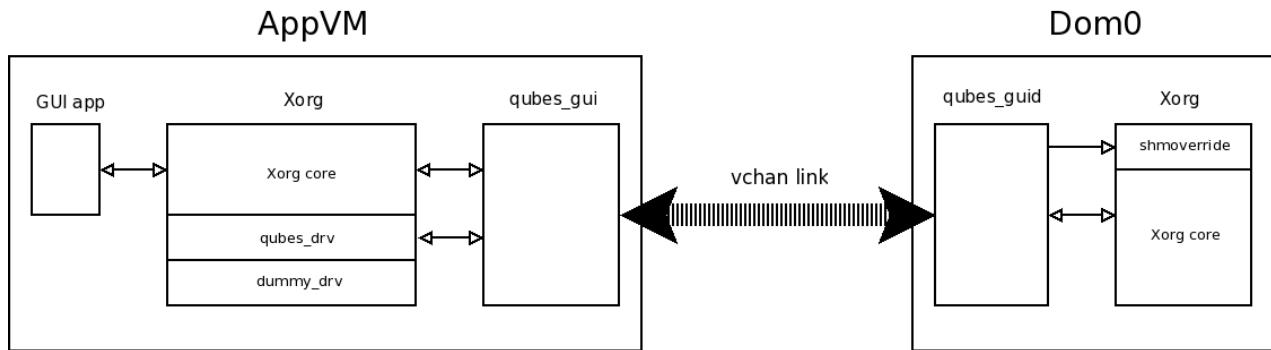
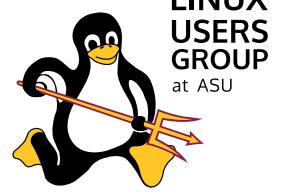
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Architecture Overview

LINUX
USERS
GROUP
at ASU



GUI Protocol Overview

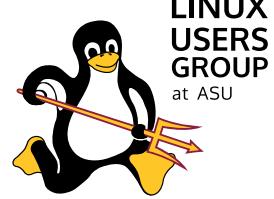


© 2022 Daylam Tayari - CC BY-NC-SA 4.0

The vchan link is the integration into Qubes of the Xen vchan protocol.

The vchan protocol is used to stream data between virtual machines without needing any locks.

GUI Drivers

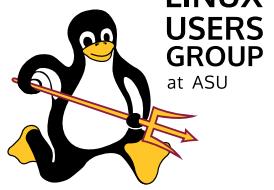


- “Hardware” Drivers
 - dummyqsb_drv – Video driver that paints onto the framebuffer in the RAM.
 - qubes_drv – Provides a virtual keyboard and mouse.
- Qubes Drivers
 - qubes_gui – Runs in the AppVM.
 - qubes_guid – Runs in dom0.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Hardware drivers aren't actually connected to any real hardware.

qubes_gui Driver

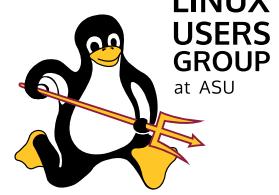


- Call XCompositeRedirectSubwindows on the root window, so that each window has its own composition buffer.
- Instruct the local Xorg server to notify it about window creation, configuration and damage events; pass information on these events to dom0.
- Receive information about keyboard and mouse events from dom0, tell qubes_drv to fake appropriate events.
- Receive information about window size/position change, apply them to the local window.

<https://www.qubes-os.org/doc/gui/> © 2022 Daylam Tayari - CC BY-NC-SA 4.0

This is all of its responsibilities.

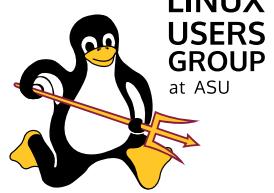
qubes_guid Driver



- Create a window in dom0 whenever an information on window creation in AppVM is received from qubes_gui.
- Whenever the local window receives XEvent, pass information on it to AppVM (particularly, mouse and keyboard data).
- Whenever AppVM signals damage event, tell local Xorg server to repaint a given window fragment.
- Receive information about window size/position change, apply them to the local window.

<https://www.qubes-os.org/doc/gui/> © 2022 Daylam Tayari - CC BY-NC-SA 4.0

Updating Windows



1. qubes_gui

1. Asks the qubes_drv driver for a list of physical memory frames that hold the composition buffer of a window.
2. Passes on this information to the qubes_guid driver in dom0 via a MFNDUMP message.

2.qubes_guid

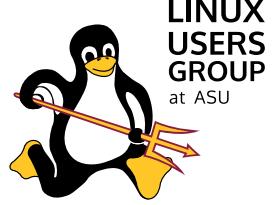
- 1.Xorg server is started with a LD_PRELOAD-ed library named shmoveoverride.so which hooks all functions calls related to shared memory.
- 2.qubes_guid creates a shared memory segment and then instructs Xorg to attach it via the MIT-SHM extension.
- 3.When Xorg tries to attach the segment via glibc shmat, shmoveoverride.so intercepts the calls and instead maps AppVM memory via xc_foreign_pagers
- 4.We can then use MIT-SHM functions to draw onto a dom0 window.

<https://www.qubes-os.org/doc/gui/> © 2022 Daylam Tayari - CC BY-NC-SA 4.0

Messages are a set list of messages that the GUI drivers use to communicate.

Full list is at the bottom of this:

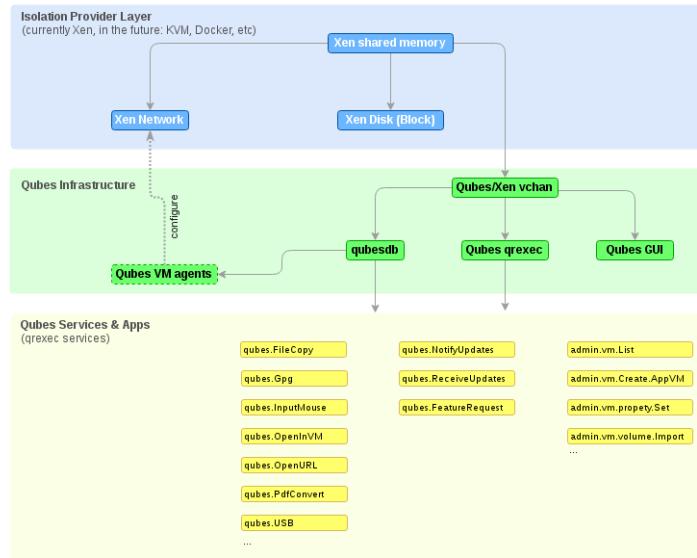
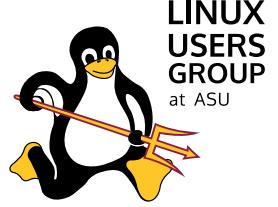
<https://www.qubes-os.org/doc/gui/>

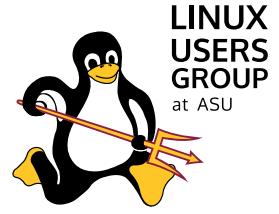


Section 4.4 – General Overview

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

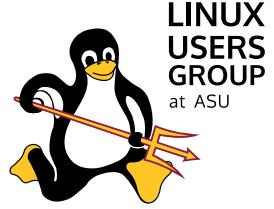
Overview





Section 5 – Security and Attack Surfaces

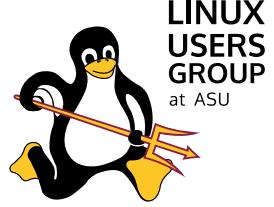
© 2022 Daylam Tayari - CC BY-NC-SA 4.0



Section 5.1 – Project Security

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

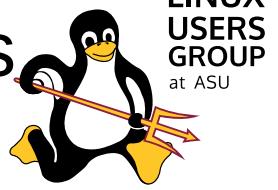
Security Team



- Handles disclosed security issues.
- Evaluates the importance of security advisories of dependencies.
- Applying security patches.
- Publishing security bulletins and canaries.
- Managing the project's PGP keys.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

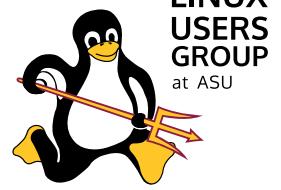
Security Team Members



- Marek Marczykowski-Górecki – Project Lead
- Simon Gaiser – Project Developer
- Joanna Rutkowska (canaries only) – Project founder

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

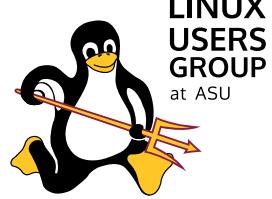
Signing and PGP Keys



- Qubes has signing keys for each major release.
- Every Qubes team member's PGP keys are published on the Qubes key site.
- Every member of the security team must sign a security release.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

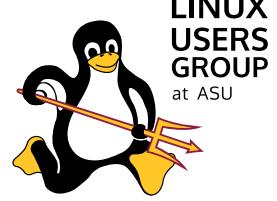
Canaries



- Canaries are periodical security announcements by the Qubes security team.
- Requires every member sign them.
- Includes specific information such as the date, its number, when the next one plans to be, etc.
- Includes proof of time by news articles and the hash of a Bitcoin block mined in the last 24 hours before the release.

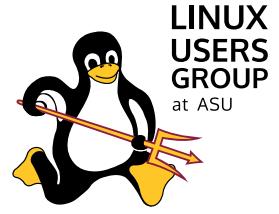
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Security Bulletins



- Special announcement issued by the security team.
- Contains summary and analysis of a specific set of vulnerabilities and their patch.

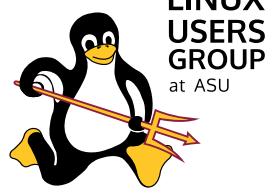
© 2022 Daylam Tayari - CC BY-NC-SA 4.0



Section 5.2 – Hardware Mitigation

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

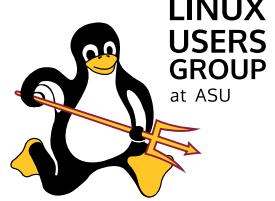
Disk Encryption and MFA



- LUKS 2 full disk encryption.
- U2F Authentication – Only an alternative
- Multi-factor Authentication – Only for restricting access to external services, not the login.

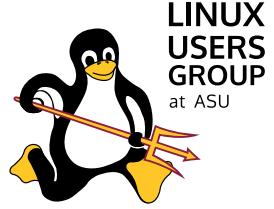
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Anti Evil Maid



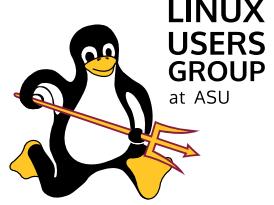
- External TPM trusted boot.
- If the whole chain is valid and properly executed, a message will be loaded validating it.
- Requires USB devices to be able to attach directly to dom0.

© 2022 Daylam Tayari - CC BY-NC-SA 4.0



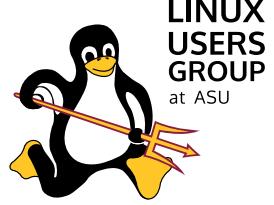
Section 5.3 – Attack Surfaces and Risk

© 2022 Daylam Tayari - CC BY-NC-SA 4.0



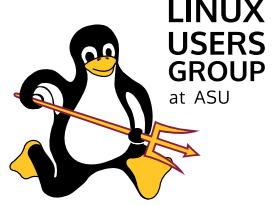
Physical Security of the Device

© 2022 Daylam Tayari - CC BY-NC-SA 4.0



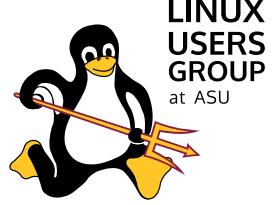
Hypervisor

© 2022 Daylam Tayari - CC BY-NC-SA 4.0



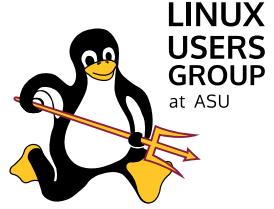
dom0 Drivers

© 2022 Daylam Tayari - CC BY-NC-SA 4.0



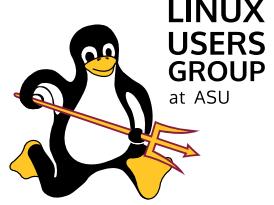
Custom-built Malware

© 2022 Daylam Tayari - CC BY-NC-SA 4.0



Real World Risk?

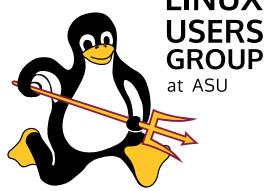
© 2022 Daylam Tayari - CC BY-NC-SA 4.0



Conclusion

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Check out the Project



- QubesOS official website: <https://www.qubes-os.org/>
- Download QubesOS:
<https://www.qubes-os.org/downloads/>
- QubesOS Forums: <https://forum.qubes-os.org/>
- QubesOS GitHub: <https://github.com/QubesOS>

© 2022 Daylam Tayari - CC BY-NC-SA 4.0

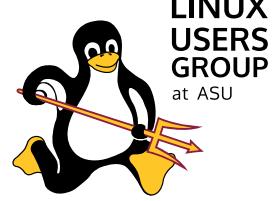
Support the Project



- QubesOS Donation:
<https://www.qubes-os.org/donate/>
- Contribute:
<https://www.qubes-os.org/doc/contributing/>

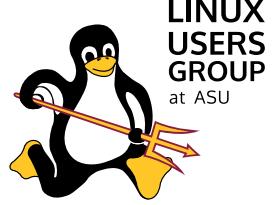
© 2022 Daylam Tayari - CC BY-NC-SA 4.0

Resources



- Qubes Website: <https://www.qubes-os.org/>
- Qubes Documentation: <https://www.qubes-os.org/doc/>
- Qubes GitHub: <https://github.com/QubesOS>
- Xen Project: <https://xenproject.org/>
- Xen Project Wiki: https://wiki.xenproject.org/wiki/Main_Page
- Invisible Things Lab: <https://invisiblethingslab.com>
- Tarsnap website and documentation: <https://www.tarsnap.com/>

© 2022 Daylam Tayari - CC BY-NC-SA 4.0



Questions?

© 2022 Daylam Tayari - CC BY-NC-SA 4.0